

Chapter 5

Class

The idea of classes goes back long before Plato. The platonic solids were classes, instances of which could be seen in the world. The platonic sphere was absolutely perfect but insubstantial. The spheres around us we could touch, but they were all imperfect in some way.

Object-oriented programming picked up on this idea by way of later western philosophers, dividing programs into classes, which are general descriptions of a whole set of similar things, and objects, which are the things themselves.

Classes are important for communication because they describe, potentially, many specific things. Class-level patterns have the largest span of any of the implementation patterns. Design patterns, by contrast, generally talk about the relationships between classes.

The following patterns appear in this chapter:

- **Class**—Use a class to say, “This data goes together and this logic goes with it.”
- **Simple Superclass Name**—Name the roots of class hierarchies with simple names drawn from the same metaphor.
- **Qualified Subclass Name**—Name subclasses to communicate the similarities and differences with a superclass.
- **Abstract Interface**—Separate the interface from the implementation.
- **Interface**—Specify an abstract interface which doesn’t change often with a Java interface.
- **Versioned Interface**—Extend interfaces safely by introducing a new sub-interface.
- **Abstract Class**—Specify an abstract interface which will likely change with an abstract class.

- Value Object—Write an object that acts like a mathematical value.
- Specialization—Clearly express the similarities and differences of related computations.
- Subclass—Express one-dimensional variation with a subclass.
- Implementor—Override a method to express a variant of a computation.
- Inner Class—Bundle locally useful code in a private class.
- Instance-specific Behavior—Vary logic by instance.
- Conditional—Vary logic by explicit conditionals.
- Delegation—Vary logic by delegating to one of several types of objects.
- Pluggable Selector—Vary logic by reflectively executing a method.
- Anonymous Inner Class—Vary logic by overriding one or two methods right in the method that is creating a new object.
- Library Class—Represent a bundle of functionality that doesn't fit into any object as a set of static methods.

Class

Data changes more frequently than logic. That is the observation that makes classes work. Each class is a declaration, “This logic goes together and changes more slowly than the data values on which it operates. These data values also go together, changing at similar rates and being operated on by related logic.” This strict separation between data-that-changes and logic-that-does-not-change isn't absolute. Sometimes the logic will be a little different based on data values; sometimes the logic varies considerably. Sometimes data doesn't change in the course of a computation. Learning how to bundle logic in classes and express variations in that logic is part of programming effectively with objects.

Organizing classes into hierarchies is a form of compression, taking the superclass and including it textually in all the subclasses. As with all compression techniques, it leaves the code more difficult to read. You have to understand the context of the superclass to be able to understand the subclass.

Using inheritance judiciously is another aspect of programming effectively with objects. Creating a subclass says, “I'm like that superclass, only different.” (Isn't it strange that we speak of *overriding* a method in a *subclass*? How much better programmers would we be if we had thoughtfully selected our metaphors?)

Classes are relatively expensive design elements in programs built from objects. A class should do something significant. Reducing the number of classes in a system is an improvement, as long as the remaining classes do not become bloated.

The patterns that follow explain how to communicate by declaring classes.

Simple Superclass Name

Finding just the right name is one of the most satisfying moments in programming. You're struggling with an idea. Oftentimes, the code has gotten complicated and it doesn't seem like it needs to be. Then, often in conversation, someone says, "Oh, I see! This is really a Scheduler." Everyone sits back and lets out a breath. The right name results in a cascade of further simplifications and improvements.

Some of the most important names to choose well are those of classes. Classes are the central anchoring concept in the design. Once classes have been named, the names of operations follow. It is rare for the reverse to be true, except if the class was poorly named in the first place.

In naming classes there is a tension between brevity and expressiveness. You'll be using class names in conversation: "Did you remember to rotate the Figure before translating it?" The names should be short and punchy. However, to make the names precise sometimes seems to require several words.

A way out of this dilemma is picking a strong metaphor for the computation. With a metaphor in mind, even single words bring with them a rich web of associations, connections, and implications. For example, in the HotDraw drawing framework, my first name for an object in a drawing was `DrawingObject`. Ward Cunningham came along with the typography metaphor: a drawing is like a printed, laid-out page. Graphical items on a page are figures, so the class became `Figure`. In the context of the metaphor, `Figure` is simultaneously shorter, richer, and more precise than `DrawingObject`.

Sometimes, good names take time to find. You may have code "done" and working for weeks, months, or (in one notable case for me) years when you discover a better name for a class. Sometimes you need to push a little harder to find a name: pull out a thesaurus, write down a list of the least-suitable names you can think of, take a walk. Sometimes you need to forge ahead with new functionality, trusting time, frustration, and your subconscious to supply a better name.

Conversation is a tool that consistently helps me find better names. Explaining the purpose of an object to another person leads me to look for rich

and evocative images to describe it. These images can lead in turn to new names.

Look for one-word names for important classes.

Qualified Subclass Name

The names of subclasses have two jobs. They need to communicate what class they are *like* and how they are *different*. Again, the balance to be struck is between length and expressiveness. Unlike the names at the roots of hierarchies, subclass names aren't used nearly as often in conversation, so they can be expressive at the cost of being concise. Prepend one or more modifiers to the superclass name to form a subclass name.

One exception to this rule is when subclassing is used strictly as an implementation sharing mechanism and the subclass is an important concept in its own right. Give subclasses that serve as the roots of hierarchies their own simple names. For example, HotDraw has a class `Handle` which presents figure-editing operations when a figure is selected. It is called, simply, `Handle` in spite of extending `Figure`. There is a whole family of handles and they most appropriately have names like `StretchyHandle` and `TransparencyHandle`. Because `Handle` is the root of its own hierarchy, it deserves a simple superclass name more than a qualified subclass name.

Another wrinkle in subclass naming is multiple-level hierarchies. Multi-level hierarchies are usually delegation waiting to happen, but while they exist they need good names. Rather than blindly prepend the modifiers to the immediate superclass, think about the name from the reader's perspective. What class does he need to know this class is like? Use that superclass as the basis for the subclass name.

Communication with people is the purpose of class names. As far as the computer is concerned, classes could simply be numbered. Class names that are too long are hard to read and format. Class names that are too short tax the reader's short-term memory. Clusters of classes whose names don't relate to each other will be difficult to comprehend and recall. Use class names to tell the story of your code.

Abstract Interface

The old adage in software development is to code to interfaces, not implementations. This is another way of suggesting that a design decision should not be visible in more places than necessary. If most of my code only knows that I am dealing with a collection I am free to change the concrete class later. However,

at some point you actually have to commit to a concrete class so the computer can perform a calculation.

By “interface” here I mean “a set of operations without implementations”. This can be represented in Java either as an interface or as a superclass. Patterns below will suggest when each is appropriate.

Every layer of interface has costs. It is one more thing to learn, understand, document, debug, organize, browse, and name. Maximizing the number of interfaces doesn’t minimize the cost of software. Pay for interfaces only where you will need the flexibility they create. Since you can’t often know in advance where you will need the flexibility of an interface, to minimize cost, combine speculating about where to introduce interfaces with adding them when flexibility is required.

Much as we complain about the inflexibility of software, there are a very large number of ways we don’t need any given system to flex. From fundamental changes like the number of bits in an integer to large-scale changes like new business models, most software doesn’t need to be flexible in most of the ways it could be.

Another economic factor in introducing interfaces is the unpredictability of software. Our industry seems addicted to the idea that if only we designed software right we would not have to change our systems. I recently read a list of reasons software changes. On the list were programmers doing a bad job of eliciting requirements, sponsors changing their mind, and on and on. The one factor that was missing from the list was legitimate change. The list assumed change was always a mistake. Why won’t one weather forecast do for all time? Because the weather changes in unpredictable ways. Why can’t we list once and for all the ways a system needs to be flexible? Because the requirements change in unpredictable ways and technology changes in unpredictable ways. This doesn’t relieve us from the responsibility to do our best to develop the system customers need right now, but it suggests that there are limits to the value of “future-proofing” software through speculation.

Putting all of these factors together—the need for flexibility, the cost of flexibility, the unpredictability of where flexibility is needed—leads me to the belief that the time to introduce flexibility is when it is definitely needed. Introducing flexibility costs because of the changes you need to make to existing software. If you can’t personally change all the software that needs to change, the costs rise further, a topic taken up in detail in the chapter on evolving frameworks.

Java’s two mechanisms for abstract interfaces, superclasses and interfaces, have different cost profiles for such changes.

Interface

One way of saying “Here’s what I want to accomplish and beyond that are details that shouldn’t concern me” is to declare a Java interface. Interfaces are one of the important innovations first put in a mass-market language in Java. Interfaces are a nice balance. They have some of the flexibility of multiple inheritance without the complexity and ambiguity. One class can declare itself as participating in multiple interfaces. Interfaces reveal only operations, not fields, so they can effectively protect users of an interface from changes in implementation.

If interfaces enable changes to their implementations, they discourage changes to the interface itself. Any addition or change to an interface requires modifying all implementors. If you can’t modify the implementations, widespread use of interfaces provides significant drag on further design evolution.

One quirk of interfaces that limits their value as a way to communicate is that all operations are required to be public. I have often wished for package-visible operations in interfaces. Making design elements a little too public isn’t a problem when they are for private use, but when publishing interfaces to a large audience it would be better to be able to be precise rather than build up inertia against future change.

Two styles of naming interfaces depend on how you are thinking of the interfaces. Interfaces as classes without implementations should be named as if they were classes (Simple Superclass Name, Qualified Subclass Name). One problem with this style of naming is that the good names are used up before you get to naming classes. An interface called `File` needs an implementation class called something like `ActualFile`, `ConcreteFile`, or (yuck!) `FileImpl` (both a suffix and an abbreviation). In general, communicating whether one is dealing with a concrete or abstract object is important, whether the abstract object is implemented as an interface or a superclass is less important. Deferring the distinction between interfaces and superclasses is well supported by this style of naming, leaving you free to change your mind later if that becomes necessary.

Sometimes, naming concrete classes simply is more important to communication than hiding the use of interfaces. In this case, prefix interface names with “I”. If the interface is called `IFile`, the class can be simply called `File`.

Abstract Class

The other way to express the distinction between abstract interface and concrete implementation in Java is to use a superclass. The superclass is abstract in

the sense that it can be replaced at runtime with any subclass, whether it is abstract in the Java sense or not.

The trade-offs for when to use an abstract class versus an interface boil down to two issues: changes to the interface and the need for a single class to support multiple interfaces simultaneously. Abstract interfaces need to support two kinds of change: change in the implementation and change of the interface itself. Java interfaces do a poor job of supporting the latter. Every change to an interface requires changes to all implementations. With a widely implemented interface, this can easily lead to paralysis of existing designs, with further evolution only available through versioned interfaces.

Abstract classes do not suffer this limitation. As long as a default implementation can be specified, new operations can be added to an abstract class without disrupting existing implementors.

One limitation of abstract classes is that implementors can only declare their allegiance to one superclass. If other views of the same class are necessary, they must be implemented by Java interfaces.

Using the keyword `abstract` with a class tells readers that they will have to do some implementation work if they want to use the class. If there is any chance to make the root of a class hierarchy useful and instantiable on its own, do so. Once on the path of abstraction, it is easy to go too far and create abstractions that never pay off. Striving to make root classes instantiable encourages you to eliminate abstractions that are unlikely to pull their weight.

Interfaces and class hierarchies are not mutually exclusive. You can provide an interface which says, “Here’s how to access this kind of functionality,” and a superclass which says, “Here’s one way to implement this functionality.” In this case, variables should be declared with the interface as their type so future maintainers are free to substitute new implementations as necessary.

Versioned Interface

What do you do when you need to change an interface but you can’t? Typically this happens when you want to add operations. Since adding an operation will break all existing implementors, you can’t do that. However, you can declare a new interface that extends the original interface and add the operation there. Users who want the new functionality use the extended interface while existing users remain oblivious to the existence of the new interface. Anywhere you want to access the new operation you must explicitly check the type of your object and downcast to the new type.

For example, consider a simple command:


```
interface Command {
    void run();
}
```

Once this interface has been published and extended a thousand times, changing it becomes expensive. However, to support undoing of commands you need a new operation. The versioned interface style of solution is this:

```
interface ReversibleCommand extends Command {
    void undo();
}
```

Existing instances of `Command` work as before. Instances of `ReversibleCommand` work anywhere a `Command` works. To use the new operation, downcast:

```
...
Command recent= ...;
if (recent instanceof ReversibleCommand) {
    ReversibleCommand downcasted= (ReversibleCommand) recent;
    downcasted.undo();
}
...
```

Using `instanceof` generally reduces flexibility by tying code to certain classes. In this case, however, it may be justified because it enables the evolution of interfaces. If you begin to have several alternative interfaces, however, clients need to do a lot of work to deal with all the variations. These are signs that it is time to rethink the design.

Alternative interfaces are an ugly solution to an ugly problem. Interfaces don't accommodate change to their structure as easily as they accommodate change to their implementations. Interfaces are likely to change, just like all design decisions. We all learn about design through implementation and maintenance. Alternative interfaces create a new programming language that is like Java but with new rules. Writing new languages is a different game with tougher rules than writing applications. However, if you are stuck in the situation of needing to extend an interface, it's nice to know how.

Value Object

While objects-with-changing-state is one valuable way to think about computation, it is not the only way to think. Mathematics has developed over millennia as a way to think about situations that can be reduced to an abstract world of absolute truth and certainty, where statements can be made about eternal verities.

Our current programming languages are a mix of the two styles. The so-called primitive types in Java belong (mostly) to the world of mathematics.

When I add 1 to a number in Java, I am making a statement of mathematics (all except that part where someone decided my computer only had to count up to 2^{32} or 2^{64} and then we should let it start over). I don't change the value of a variable when I add 1: I create a new value. There is no way to change 0, as you can with most objects.

This functional style of computing never changes any state, it only creates new values. When you have a (perhaps momentarily) static situation about which you'd like to make statements or about which you'd like to ask questions, then the functional style is appropriate. When the situation is changing over time, then state is appropriate. Some situations could be thought of in either way. How can you tell which way is most helpful?

For example, you can represent drawing a picture as changes to the state of some graphics medium like a bitmap. Alternatively, you can describe the same picture with a static description (Figure 5.1).




Figure 5.1 Graphics represented as procedures and objects

Which of these representations is most useful depends to some extent on personal preference, but it also depends on the complexity of the pictures to be drawn and how often they change.

Procedural interfaces are more common than functional interfaces. One problem with procedural interfaces is that the sequence of procedure calls becomes an important (but often implicit) part of the meaning of the interface. Changing such a program is touchy and difficult because seemingly small changes have unintended consequences when the implicit meaning of the sequence is changed.

The beauty of mathematical representations is that sequence seldom matters. You are creating a world in which you can make absolute, timeless statements. Create micro-worlds of mathematics wherever possible. Manage them from an object with changing state.

For example, implement an accounting system by making the basic transactions unchanging mathematical values.




Figure 5.2 *State-changing objects referring to immutable objects at the edges*

```
class Transaction {
    int value;
    Transaction(int value, Account credit, Account debit) {
        this.value= value;
        credit.addCredit(this);
        debit.addDebit(this);
    }
    int getValue() {
        return value;
    }
}
```

There is no way to change any of the values of a Transaction once it has been created. What's more, the constructor makes the statement that all transactions are posted to two accounts. When I read this code, I know that I don't have to worry about transactions floating around loose or about transactions having their values changed after they have been posted.

To implement value-style objects (that is, objects that act like integers rather than like holders of changing state), first draw the boundary between the world of state and the world of value. In the example above, a Transaction is a value, an Account holds changing state. Set all state in a value-style object in the constructor, with no other field assignments elsewhere in the object. Operations on value-style objects always return new objects. These objects must be stored by the requestor of the operation.

```
bounds.translateBy(10, 20); // mutable Rectangle
bounds= bounds.translateBy(10, 20); // value-style Rectangle
```

The biggest argument against value-style objects has always been performance. The need to create all those intermediate objects can put a strain on the memory management system. In the overall cost of programming, this argument doesn't often hold up because most parts of the program are not

performance bottlenecks. Other reasons not to use value-style objects are unfamiliarity with the style and difficulty drawing boundaries between parts of the system where state changes and parts of the system where objects don't change. Objects that are mostly value-style are the worst of both worlds, since the interfaces tend to be more complicated but you can't safely make assumptions about state not changing.

Having gotten this far, I sense that there is much more to be written about programming with the three main styles—objects, functions, and procedures—and how to blend them effectively. For purposes of this book, I'll close by reiterating that sometimes your programs will best be expressed as a combination of state-changing objects and objects representing mathematical values.

Specialization

Communicating the interplay between the similarities and differences of computations makes your programs easier to read, use, and modify. In practice, each program is not unique. Many express similar ideas, and, often, many parts of the same program express similar ideas. Expressing similarities and differences clearly allows readers to understand the existing code, discover if their current intentions are covered by one of the existing variations, or, if not, how to best either specialize the existing code to their needs or write entirely new code.

The simplest variations are those where state differs. The string “abc” is different from “def”. The algorithms that operate on the two strings are identical. For example, the length of all strings is calculated in the same way.

The most complex variations are total differences in logic. A symbolic integration routine has no logic in common with a mathematical typesetting routine, even though the two might share exactly the same input.

In between these two extremes—identical logic with different data and different logic with identical data—lies the huge common ground of programming. Data can be mostly the same but a little different. Logic can be mostly the same but a little different. (I would guess that the symbolic integration routine and the mathematical typesetting routine share little code.) Even the line between logic and data is blurry. A flag is boolean data but it affects the flow of control. A helper object can be stored in a field but used to affect a computation.

The patterns that follow are a variety of techniques to communicate similarity and difference, primarily in logic. Variations in data don't seem as

complicated or subtle. Effective expressions of similarity and difference in logic open up new opportunities for further expansion of the code.

Subclass

Declaring a subclass is a way of saying, “These objects are like those except...” If you have the right superclass, creating a subclass can be a powerful way to program. With the right method to override, you can introduce a variant of an existing calculation with a few lines of code.

When objects first became popular, subclassing seemed like a magic pill. First, subclasses were used for classification—a Train was a subclass of Vehicle regardless of whether they shared any implementation. In time, some people saw that since what inheritance did was share implementation, it could most effectively be used to factor out common bits of implementation. Quickly, though, the limitations of subclassing became apparent. First, it’s a card you can only play once. If you discover that some set of variations isn’t well expressed as subclasses, you have some work to do to disentangle the code before you can restructure it. Second, you have to understand the superclass before you can understand the subclass. As the superclasses become more complicated this becomes more of a limitation. Third, changes to a superclass are risky, since subclasses can rely on subtle properties of the superclass’s implementation. Finally, all of these problems are compounded by deep inheritance hierarchies.

A particularly pernicious use of inheritance is creating parallel hierarchies, where for each subclass in *this* hierarchy you need a subclass in *that* hierarchy. This is a form of duplication, creating implicit coupling between the class hierarchies. To successfully introduce a new variation you need to change both hierarchies. While I often see parallel hierarchies that I can’t immediately figure out how to eliminate, the effort to do so improves the design.




Figure 5.3 Parallel hierarchies

One example of this was an insurance system (Figure 5.3). Something is definitely wrong with this picture, because an `InsuranceContract` cannot refer to a `PensionProduct` nor is it attractive to move the `product` field down to the subclasses. The solution, which we never reached but took a year to come close to, was to move the variation around so `Contract` worked the same whether it was used for insurance or a pension. This required creating a new object to represent the prospective cash flows (Figure 5.4).




Figure 5.4 *Hierarchy with duplication eliminated*

With all these warnings in mind, subclassing can be a powerful tool for expressing theme-and-variations computations. The right subclass can help many people to express exactly the computation they want with a method or two. One key to achieving useful subclasses is to thoroughly factor the logic in the superclass into methods that do one job. When writing a subclass, you should be able to override exactly one method. If the superclass methods are too big, you'll have to copy code and edit it (Figure 5.5).




Figure 5.5 *Code copied and modified in a subclass*

Copied code introduces an ugly implicit coupling between the two classes. You can't safely change the code in the superclass without examining and potentially changing all the places to which it has been copied.

My goal in design is to be able to switch between strategies at will, depending on the needs of the code as it now exists. Visualize the code expressed with conditionals, with subclasses, with delegation. Does it seem like there are advantages to a different strategy than the one you are currently using? Take a few steps in that direction and see if you improve the code.

A final limitation of subclassing is that it cannot be used to express changing logic. The variation you want must be known when you create an object and can't be changed thereafter. You'll need to use conditionals or delegation to express logic that changes.

Implementor

The polymorphic message is the fundamental way to express a choice in a program built from objects. For the message to do its work of choosing, there needs to be more than one kind of object to potentially receive the message.

Implementing the same protocol multiple times, whether expressed with a Java interface and an `implements` declaration or as a subclass expressed with `extends`, is a way of saying, "From the point of view of one part of the calculation, as long as something happens matching the intention of the code, the details of what happens are irrelevant."

The beauty of polymorphic messages is that they open a system to variation. If a part of the program writes some bytes to another system, the introduction of an abstract `Socket` enables the implementation of the socket to vary without affecting the calling code. Compared to the procedural expression of the same intention, with its explicit and closed conditional logic, the object/message version is clearer, separating the expression of the intention (write some bytes) from the implementation (call a TCP/IP stack with certain parameters). At the same time, expressing the computation as objects and messages opens the system to future variation undreamt of by the original programmers. This fortuitous combination of clarity of expression and flexibility is why object languages have become the dominant programming paradigm.

This supreme resource is easy to squander by writing procedural programs in Java. The patterns here are intended to help you express logic that is both clear and extendable.

Inner Class

Sometimes you need to package part of a computation but you don't want to incur the cost of a whole new class with its own file. Declaring small, private

classes (inner classes) gives you a low-cost way of having many of the benefits of a class without all of its costs.

Sometimes an inner class extends only `Object`. Some inner classes extend another superclass, which is useful for expressing refinements of other classes that are only locally interesting.

One of the features of inner classes is that when their instances are created, they are secretly passed a copy of the object that is creating them. This is handy when you want to access the enclosing instance's data without making the relationship between the two classes explicit:

```
public class InnerClassExample {
    private String field;

    public class Inner {
        public String example() {
            return field; // Uses the field from the enclosing instance
        }
    }

    @Test public void passes() {
        field= "abc";
        Inner bar= new Inner();
        assertEquals("abc", bar.example());
    }
}
```

However, in the inner class above, there is not really a no-arg constructor, even if you declare one. This is a problem when creating instances of inner classes by reflection.

```
public class InnerClassExample {
    public class Inner {
        public Inner() {
        }
    }
}

@Test(expected=NoSuchMethodException.class)
public void innerHasNoNoArgConstructor() throws Exception {
    Inner.class.getConstructor(new Class[0]);
}
}
```

To get an inner class that is completely detached from its enclosing instances, declare it static.

Instance-Specific Behavior

In theory, all instances of a class share the same logic. Relaxing this constraint enables new styles of expression. All of these styles, though, come at a cost. When the logic of an object is completely determined by its class, readers can read the code in the class to see what is going to happen. Once you have instances with different behavior, you have to look at live examples or analyze the data flow to understand how a particular object is going to behave.

Another step up in the cost of instance-specific behavior is when the logic changes as the computation progresses. For ease of code reading, try to set instance-specific behavior when an object is created and don't change it afterward.

Conditional

If/then and switch statements are the simplest form of instance-specific behavior. Using conditionals, different objects will execute different logic based on their data. Conditionals as a form of expression have the advantage that the logic is still all in one class. Readers don't have to go navigating around to find the possible paths for a computation. However, conditionals have the disadvantage that they can't be modified except by modifying the code of the object in question.

Each path of execution through a program has some probability of being correct. Assuming that the probabilities of correctness for the paths are independent, the more paths through a program the less likely the program is to be correct. The probabilities aren't entirely independent, but they are independent enough that programs with many paths are more likely to have defects than those with few paths. The proliferation of conditionals reduces reliability.

This problem is compounded when conditionals are duplicated. Consider a simple graphic editor. The figures will need a `display()` method:

```
public void display() {  
    switch (getType()) {  
        case RECTANGLE :  
            //...  
            break;  
        case OVAL :  
            //...  
            break;  
        case TEXT :  
            //...  
            break;  
        default :
```

```

        break;
    }
}

```

Figures will also need a method to determine whether a point is contained within them:

```

public boolean contains(Point p) {
    switch (getType()) {
        case RECTANGLE :
            //...
            break;
        case OVAL :
            //...
            break;
        case TEXT :
            //...
            break;
        default :
            break;
    }
}

```

Suppose now you want to add a new kind of figure. First, you must add a clause to every switch statement. Second, to make this change you have to modify the Figure class, putting all of the existing functionality at risk. Lastly, everyone who wants to add new figures must coordinate their changes of a single class.

These problems can all be eliminated by converting the conditional logic to messages, either with subclasses or delegation (which technique serves best depends on the code). Duplicate conditional logic or logic where the processing is very different based on which branch of a conditional is taken is generally better expressed as messages instead of explicit logic. Also, conditional logic that changes frequently is better expressed as messages to simplify changing one branch while minimizing the effects on other branches.




Figure 5.6 Conditional logic represented by subclasses and delegation

In short, the strengths of conditionals—that they are simple and local—become liabilities when they are used too widely.

Delegation

Another way to execute different logic in different instances is to delegate work to one of several possible kinds of objects. The common logic is held in the referring class, the variations in the delegates.

An example of using a delegate to capture variation is handling user input in a graphical editor. Sometimes a button press means “create a rectangle”, sometimes it means “move a figure”, and so on.

One way to express the variation between the tools is with conditional logic:

```
public void mouseDown() {
    switch (getTool()) {
        case SELECTING :
            //...
            break;
        case CREATING_RECTANGLE :
            //...
            break;
        case EDITING_TEXT :
            //...
            break;
        default :
            break;
    }
}
```

This has all the problems of conditionals discussed above: adding a new tool requires modifying the code and the duplication of the conditional (in `mouseUp()`, `mouseMove()`, etc.) makes adding new tools complicated.

Subclassing is not an immediate answer either because the editor needs to change tools during its lifetime. Delegation allows that flexibility.

```
public void mouseDown() {
    getTool().mouseDown();
}
```

The code that used to live in the clauses of the switch statement is moved to the various tools. Now new tools can be introduced without modifying the code of the editor or the existing tools. Reading the code requires more navigation, however, because the mouse-down logic is spread over several classes. Understanding how the editor will behave in a given situation requires that you understand what kind of tool it is currently using.

Delegates can be stored in fields (a “pluggable object”), but they can also be computed on the fly. JUnit 4 dynamically computes the object that will run the tests in a given class. If a class contains old-style tests, one delegate is created, but if the class contains new-style tests, a different delegate is created. This is a mix of conditional logic (to create the delegates) and delegation.

Delegation can be used for code sharing as well as instance-specific behavior. An object that delegates to a `Stream` may be involved in instance-specific behavior, if the type of `Stream` can change at runtime, or it may be sharing the implementation of `Stream` with all the other users.

A common twist on delegation is to pass the delegator as a parameter to a delegated method.

```
GraphicsEditor
public void mouseDown() {
    tool.mouseDown(this);
}

RectangleTool
public void mouseDown(GraphicsEditor editor) {
    editor.add(new RectangleFigure());
}
```

If a delegate needs to send a message to itself, “itself” is ambiguous. Sometimes the message should be sent to the delegating object. Sometimes the message should be sent to the delegate. In the example below the `RectangleTool` adds a figure, but to the delegating `GraphicsEditor`, not to itself. The `GraphicsEditor` could have been passed as a parameter to the delegated `mouseDown()` method, but in this case it seemed simpler to store a permanent back-reference in the tool. Passing the `GraphicsEditor` as a parameter makes it possible to use the same tool in multiple editors, but if this isn’t important the code with the backpointer may be simpler.

```
GraphicsEditor
public void mouseDown() {
    tool.mouseDown();
}

RectangleTool
private GraphicsEditor editor;
public RectangleTool(GraphicsEditor editor) {
    this.editor= editor;
}
public void mouseDown() {
    editor.add(new RectangleFigure());
}
```

Pluggable Selector

Let's say you need instance-specific behavior, but only for one or two methods, and you don't mind having all the variants of the code live in one class. In this case, store the name of the method to be invoked in a field and invoke the method by reflection.

Originally, each test in JUnit had to be stored in its own class (Figure 5.7). Each subclass only had one method. Classes seemed conceptually heavy as a way to represent a single class.




Figure 5.7 Trivial subclasses to represent different tests

By implementing a generic `runTest()`, `ListTests` with different names run different test methods. The name of the test is assumed to also be the name of a method which is retrieved and run when the test is run. Here is the simple version of the code to implement the pluggable selector version of running a test.

```

String name;
public void runTest() throws Exception {
    Class[] noArguments= new Class[0];
    Method method= getClass().getMethod(name, noArguments);
    method.invoke(this, new Object[0]);
}

```

The simplified class hierarchy uses a single class (Figure 5.8). As with all code compression techniques, the modified code is only easy to read if you understand the “trick”.

When pluggable selectors first became widely known, people tended to overuse them. You would be looking at some code, decide it couldn't possibly be called, delete, and have the system break because it was invoked by a pluggable selector somewhere. The costs of using pluggable selectors are considerable, but a limited use to solve a difficult problem may justify the cost.

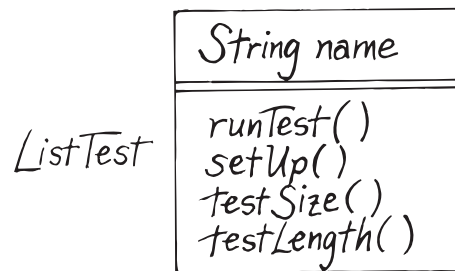


Figure 5.8 Pluggable selector helps pack tests into a single class

Anonymous Inner Class

Java offers one more alternative for instance-specific behavior, anonymous inner classes. The idea is to create a class that is only used in one place, that can override one or more methods for strictly local purposes. Because it is only used in one place, the class can be referred to implicitly instead of by name.

Effective use of anonymous inner classes relies on having an extremely simple API—like implementing `Runnable` with its one method `run()`—or having a superclass that provides most of the needed implementation so the anonymous inner class can be implemented simply. The code for the anonymous inner class interrupts the presentation of the code in which it is embedded, so it needs to be short so it doesn't distract the reader.

Anonymous inner classes have the limitations that the code to be set in the instance must be known when you write the class (unlike delegates, which can be added later) and it cannot be changed once an instance has been created. Anonymous inner classes are difficult to test directly and so shouldn't contain complicated logic. Because they are un-named, you don't have the opportunity to express your intention for an anonymous inner class with a well-chosen name.

Library Class

Where do you put functionality that doesn't fit into any object? One solution is to create static methods on an otherwise-empty class. No one is expected to ever create instances of this class. It is just there as a holder for the functions in the library.

While library classes are fairly common, they don't scale well. Putting all the logic into static methods forfeits the biggest advantage of programming with objects: a private namespace of shared data that can be used to help simplify logic. Try to turn library classes into objects whenever possible.

Sometimes this is as simple as finding a better home for a method. The `Collections` library class, for example, has a method `sort(List)`. Such a specific parameter is a hint that this method probably belongs on `List` instead.

An incremental way to convert a library class to an object is to convert the static methods to instance methods. Maintain the same interface at first by having the static method delegate to an instance method. In a class called `Library`, for example,

```
public static void method(...params...) {  
    ...some logic...  
}
```

becomes:

```
public static void method(...params...) {  
    new Library().instanceMethod(...params...);  
}  
private void instanceMethod(...params...) {  
    ...some logic...  
}
```

Now, if several of the methods have similar parameter lists (and if they don't then the methods probably belong in different classes), convert the method parameters to constructor parameters:

```
public static void method(...params...) {  
    new Library(...params...).instanceMethod();  
}  
private void instanceMethod() {  
    ...some logic...  
}
```

Next change the interface by moving instance creation to the clients and eliminating the static methods.

```
public void instanceMethod(...params...) {  
    ...some logic...  
}
```

This experience may give you ideas for how to rename the class and methods so client code reads clearly.

Conclusion

A class bundles together related state. The next chapter presents patterns that communicate decisions about state.

Chapter 6

State

The patterns in this chapter describe how to communicate your use of state. Objects are convenient packages of *behavior* which is presented to the outside world and *state* which is used to support that behavior. One of the advantages of objects is that they mince all of the state of a program into tiny pieces, each effectively its own little computer. Large libraries of state, promiscuously referenced, make further changes to code difficult because the effect of a code change on the state is hard to predict. With objects, it is easier to analyze what state will be affected by a change, because the namespace of referenceable state is so much smaller.

The chapter contains the following patterns:

- State—Compute with values that change over time.
- Access—Maintain flexibility by limiting access to state.
- Direct Access—Directly access state inside an object.
- Indirect Access—Access state through a method to provide greater flexibility.
- Common State—Store the state common to all objects of a class as fields.
- Variable State—Store state whose presence differs from instance to instance as a map.
- Extrinsic State—Store special-purpose state associated with an object in a map held by the user of that state.
- Variable—Variables provide a namespace for accessing state.
- Local Variable—Local variables hold state for a single scope.
- Field—Fields store state for the life of an object.

- **Parameter**—Parameters communicate state during the activation of a single method.
- **Collecting Parameter**—Pass a parameter to collect complicated results from multiple methods.
- **Parameter Object**—Consolidate frequently used long parameter lists into an object.
- **Constant**—Store state that doesn't vary as a constant.
- **Role-Suggesting Name**—Name variables after the role they play in a computation.
- **Declared Type**—Declare a general type for variables.
- **Initialization**—Initialize variables declaratively as much as possible.
- **Eager Initialization**—Initialize fields at instance creation time.
- **Lazy Initialization**—Initialize fields whose values are expensive to calculate just before they are first used.

State

The world persists. If a minute ago the sun was high in the sky, you can be sure it is still high in the sky, but moved a bit. If I cared to calculate, I could predict its new position based on my previous observation, knowledge of the earth's rotation, and measurement of the passage of time.

Thinking of the world as things that change has proven useful for a long time. The Native Americans in the area where I live watched Mt. McLaughlin in the spring. When the snow melted sufficiently that the outline of a flying eagle appeared in the remaining snow, it was time to move down to the Rogue River to catch the spring run of salmon. The state of the snow on the mountain was a valuable clue to the presence of a yummy meal swimming in the water some distance away.

When computing pioneers picked metaphors for programming computers, they latched onto this idea of state changing over time. The human brain has a wide variety of strategies, built-in and learned, for dealing with state.

However, state also poses problems for programmers. As soon as you assume what some bit of state is, your code is at risk. You might assume incorrectly or the state might change. Many desirable programming tools, like automated refactorers, are easier to construct if there is no notion of state.

Finally, concurrency and state don't play well together. Many of the problems of parallel programs vanish if there is no state.

Functional programming languages dispense altogether with changing state. None of these has ever become popular. I think state is a valuable metaphor for us since our brains are structured and conditioned to deal with changing state. Single assignment or variableless programming forces us to discard too many effective thinking strategies to be an attractive choice.

Object languages are a coping strategy for dealing with state. They offer the opportunity to avoid the problem of state changing “behind your back” by partitioning the state in the system into discrete little chunks, each of which has strictly limited access to the others. It is easier to keep track of a handful of bytes than mega- or gigabytes. The problem of incorrectly assuming the value of some state remains, but with objects you have the chance to quickly and accurately review all access to a variable.

A key to managing state effectively is putting similar state together and making sure different state stays apart. Two clues that two bits of state are similar are if the two bits are used in the same computation and if the two bits live and die at the same time. If two pieces of state are used together and have the same lifetime, then having them close to each other is probably a good idea.

Access

One dichotomy in programming languages is the distinction between accessing stored values and invoking computations. The two concepts are understandable in terms of each other. Accessing memory is like invoking a function that returns the currently stored values. Invoking a function is like reading a memory location, the contents of which just happen to be computed and not simply returned. Nevertheless, our programming languages separate invoking computation and accessing memory, so we need to be able to communicate effectively about the difference.

Deciding what to store and what to compute affects the readability, flexibility, and performance of programs. Sometimes these goals conflict with each other and with your programming preference. Sometimes the context changes so that yesterday's reasonable partition between stored and computed no longer makes sense. Making workable decisions today and maintaining the flexibility to change your mind in the future is a key to good software development. It is that need for future change that makes it important for you to communicate your store-versus-compute decisions clearly.

One of the goals of objects was to manage storage. Each object acts as its own little computer with its own memory, isolated to some degree from the

other little computers. Current languages, including Java, blur the boundaries between objects by offering public fields. Ease of inter-object access isn't worth the loss of independence between objects.

Direct Access

The simplest way to say “I’m fetching data” or “I’m storing data” is to use direct variable access:

```
x= 10;
```

Direct variable access has the advantage of clarity of expression. When I read `x= 10;`, I know exactly what is going to happen. This clarity comes with a loss of flexibility. If I store a value in a variable, that's all I can do. If I store values into that variable from many parts of the program, then to make a change I will likely have to change all those parts of the program.

The other downside of direct access is that it is an implementation detail, below the level of most of my thoughts while programming. Setting a variable to 1 may cause my garage door to open, but code that reflects this implementation detail won't communicate well. Compare:

```
doorRegister= 1;
```

with:

```
openDoor();
```

or, with objects:

```
door.open();
```

Most of the thoughts I think while programming have nothing to do with storage. Widespread direct access clutters communication. For those parts of the program where I really am thinking about what is stored where, I use direct access to communicate those thoughts. Storage decisions play a different role for different programmers, so there is no one policy for using direct access that will fit everyone. People keep trying to formulate such rules: direct access only inside accessor methods, and maybe inside constructors too; direct access only inside a single class or inside a class and all its subclasses, or maybe inside an entire package. There is no universal rule. Programmers need to think, communicate, and learn. That's part of being professional.

Indirect Access

You can hide accesses and changes to state behind method invocations. These accessor methods provide flexibility at the cost of clarity and directness. Clients no longer assume that a certain value is stored directly. Thus, you are able to change your mind about storage decisions without affecting client code.

My default strategy for accessing state is to allow direct access inside of a class (including inner classes) and indirect access for clients. This strategy has the advantage that it allows clear, direct access to state for most accesses. Note: if most accesses to an object's state are outside the object, there is a deeper design problem lurking.

Another strategy is to use indirect access exclusively. I find this results in a loss of clarity. Most getting and setting methods are trivial. They often outnumber methods that perform useful work, making the code hard to read. All those getting and setting methods are mighty tempting too. Rather than figure out where a calculation belongs, it is often expedient to implement it wherever and use the accessor methods to deliver the necessary state to make it work.

One definite case for indirect access is where two pieces of data are coupled. Sometimes this coupling is very direct, as in a cached value:

```
Rectangle void setWidth(int width) {  
    this.width= width;  
    area= width * height;  
}
```

Other times the coupling is less direct, through a listener:

```
Widget void setBorder(int width) {  
    this.width= width;  
    notifyListeners();  
}
```

Such coupling is unattractive (it is easy to forget to maintain the implied constraints) but may be the best available option. In such a case, indirect access is best.

Common State

Many calculations share the same data elements even if the values are different. When you find such a calculation, communicate it by declaring fields in a class. For example, all calculations with cartesian points require an abscissa and ordinate. Since all cartesian points share the need for these values, they are most clearly expressed as fields:

```
class Point {  
    int x;  
    int y;  
}
```

Contrast this technique with variable state, where objects of the same class potentially have different data elements. The advantage of common state is that it is clear from reading the code, either the fields themselves or the complete constructor, what data is necessary to have a well-formed object. Your reader will want to know what it takes to successfully invoke the functionality in your object. Common state communicates this clearly and precisely.

Common state in an object should all have the same scope and lifetime. Sometimes I am tempted to introduce a field that is only used by a subset of the methods in an object, or that is only valid while one method is being computed. In such cases I can invariably improve my code by finding somewhere else to store the data in question, perhaps a parameter or a helper object.

Variable State

Sometimes the same object needs different data elements depending on how it is used. It's not just the values that differ; whole different elements are present in objects of the same class.

Variable state is often stored as a map whose keys are the names of the elements (represented as strings or enumerations) and whose values are the data values.

```
class FlexibleObject {  
    Map<String, Object> properties= new HashMap<String, Object>();  
    Object getProperty(String key) {  
        return properties.get(key);  
    }  
    void setProperty(String key, Object value) {  
        properties.set(key, value);  
    }  
}
```

Variable state is much more flexible than common state. Its primary failing is that it doesn't communicate well. What data elements need to be present for an object with only variable state to function correctly? Only a careful reading of the code and perhaps watching execution will help you answer this question.

I have read code where the programmer overused variable state. Every object of a given class had exactly the same keys in its property map. It would have been much easier for me to read had the same information appeared as field declarations.

One case where variable state seems justified is where the state of one field implies the need for other fields. For example, if I have a widget whose bordered flag is true, then I can also have `borderWidth` and `borderColor`. I could communicate this with variable state, as in the design on the top of Figure 6.1.

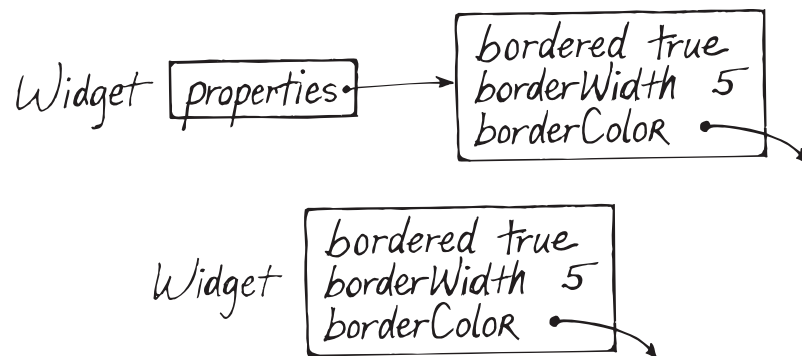


Figure 6.1 Border represented by variable state and common state

Common state can also communicate this, as in the bottom of Figure 6.1.

The common state solution violates the principle that all variables in an object should have the same lifetime. Polymorphism provides a clearer explanation of the situation. One class represents the unbordered state and another the bordered state. Bordered has common state to represent its parameters.




Figure 6.2 A helper object cleans up the design

The presence of several variables that share a common prefix is a clue that a helper object of some sort may be useful.

Use common state wherever possible. Use variable state for the fields in an object that may or may not be needed depending on usage.

Extrinsic State

Sometimes a part of your program needs state associated with an object, but the rest of the system doesn't care. For example, information about where an object is stored on disk is useful for the persistence mechanism but not the rest of the code. Putting this data in a field would violate the principle of symmetry. All the rest of the fields are useful for the whole system.

Store special-purpose information associated with an object near where it will be used instead of in the object. In the example above, the persistence mechanism would store an `IdentityMap` whose keys are the objects stored and whose values are the information about where they are stored.

One weakness of extrinsic state is that it makes copying an object difficult. Replicating an object with extrinsic state isn't as simple as replicating its fields. Instead, all the extrinsic state must also be copied correctly, which could require different handling depending on how that state is used. Another weakness is the difficulty in debugging objects with extrinsic state. A conventional inspector doesn't show all the data associated with an object. Because of these difficulties, extrinsic state is rare but useful when necessary.

Variable

In Java, objects are referred to by variables. Readers need to know about the scope, lifetime, role, and runtime type of variables. While elaborate schemes have been invented to communicate all this information in variable names, simplifying the code using simple names is preferable.

The scope of variables, the extent within which they can be referenced, can be of three types: locals, accessible only within the current scope; fields, which can be accessed anywhere within an object; or static, which can be accessed by any object of a given class. Scope of fields can be extended by the modifiers `public`, `package` (the default, and an odd choice for a default since it is the least used), `protected`, and `private`.

If you make liberal use of all the combinations available, then it is important to readers that you make the distinctions clear at the point of reference by encoding them in the name. However, to reduce coupling you should mostly use locals and fields with only an occasional static field and the `private` modifier. By using this limited set of the possible combinations, context suffices to tell the reader whether he is looking at a local or a field. If he can see the declaration the variable is a local, and if he can't then it is a field. This eliminates the need to include scope information in the variable's name. In return you have code with uniform, easily read variable names. All of this presupposes that you can

break your code up into small chunks, a property that you can accomplish by applying the other implementation patterns, most notably composed methods.

The lifetime of variables can be smaller than their scope. A field could only be valid while a certain method is active on the stack. That would be ugly. Work hard to ensure that the lifetime of variables is close to their scope. Additionally, make sure that sibling variables (those defined in the same scope) all have the same lifetime.

The type of a variable is adequately communicated through the type declaration. Make sure the declared type communicates as clearly as possible (see “Declared Type”). The one exception to this advice is that the names of variables that hold multiple values (those containing a collection) should be plural. The difference between a single value and multiple values is important for readers.

With scope, lifetime, and type adequately communicated in other ways, the name can be used to convey the role of the variable in the computation. By reducing the information to be conveyed to a minimum, you are free to choose simple names that read well.

Local Variable

Local variables are only accessible from their point of declaration to the end of their scope. Following the principle that information should spread as little as possible, declare local variables just before they are used and in the innermost possible scope.

There are a handful of common roles for local variables:

- **Collector:** a variable that collects information for later use. Often the contents of collectors are returned as the value of a function. When a collector will be returned, name it `result` or `results`.
- **Count:** a special collector that collects the count of some other objects.
- **Explaining:** if you have a complicated expression, assigning bits of the expression to local variables can help readers navigate the complexity:

```
int top= ...;
int left= ...
int height= ...;
int bottom= ...;
return new Rectangle(top, left, height, width);
```

While not computationally necessary, the explaining local variables help what would otherwise be a long, complicated expression.

Explaining locals are often a step towards helper methods. The expression becomes the body of the method, and the name of the local variable suggests a name for the method. Sometimes these helpers are introduced to simplify the calling method. Sometimes they help eliminate the duplication of common expressions.

- **Reuse:** when an expression's value changes but you need to use the same value more than once, store the value in a local variable. For example, if you need the same timestamp for several objects, you can't fetch the time fresh for each object:

```
for (Clock each: getClocks())
    each.setTime(System.currentTimeMillis());
```

Instead, a reusing local variable freezes time for your purposes:

```
long now= System.currentTimeMillis();
for (Clock each: getClocks())
    each.setTime(now);
```

- **Element:** the final common use of local variables is to hold the elements of a collection that is being iterated. As in the example above, `each` is a clear, simple name for an element local variable. If I want to know “each what?” I can glance at the `for` statement above.

For nested loops, append the name of the collection to the element local name to distinguish them:

```
broadcast() {
    for (Source eachSender: getSenders())
        for (Destination eachReceiver: getReceivers())
            ...;
}
```

Field

The scope and lifetime of a field is the same as the object to which it is attached. Because the primary allegiance of fields is to the object as a whole, declare fields together either at the beginning or end of the class. At the beginning, the declarations give a reader important context to use while reading the rest of the code. Leaving the declarations until last sends the message, “Behavior is king; data is an implementation detail.” While I agree philosophically with the statement that logic is more important than data in object programs, I still like to read the declarations first, wherever they are, when I am reading code.

One of your options with a field is to declare it `final`. This tells readers that the value of the field will not change after the constructor has run. While I

mentally keep track of which of my fields are final and which are not, I don't declare fields explicitly. The extra clarity is not worth the extra complexity for me, but if I were writing code that would be changed by many people over a long time, it seems worth making the distinction between final and volatile fields explicit.

The list of roles for fields is not as comprehensive as the list for local variables. However, here are a few common roles for fields:

- **Helper:** helper fields hold references to objects used by many of an object's methods. If an object is passed as a parameter to many methods, consider replacing the parameter with a helper field set in the complete constructor.
- **Flag:** boolean flag fields say, "This object can act in two different ways." If there is a setter method for the flag, it additionally says, "... and the behavior can change during the life of the object." Flag fields are fine if they are only used in a few conditionals. If the code making decisions on the basis of the flag is duplicated, consider changing to a strategy field instead.
- **Strategy:** when you want to express that there are alternative ways to perform some part of an object's computation, store in a field an object performing just the variable part of the computation. If this variation in behavior doesn't change during an object's lifetime, set the strategy field in the complete constructor. Otherwise, provide methods for changing it.
- **State:** state fields are like strategy fields in that part of the behavior of the object is delegated to them. However, state fields, when triggered, set the following state themselves. Strategy fields are changed, if at all, by other objects. State machines implemented this way can be difficult to read, since the states and transitions are not expressed in one place. However, for simple state machines they can suffice.
- **Components:** these fields hold objects or data "owned" by the referring object.

Parameter

Besides non-private variables (fields or static fields), the only way to communicate state from one object to another is through parameters. Because non-private variables introduce strong coupling between classes and because that coupling tends to grow over time, parameters are preferable in all cases where both static fields and parameters are possible.

The coupling introduced by a parameter is weaker than the coupling introduced by a permanent reference from one object to another. For example, calculations in the interior of a tree structure sometimes need the parent of a node. Rather than have a permanent reference to the parent (Figure 6.3), passing a parameter to those methods requiring it weakens the coupling between the nodes. With no permanent reference to a parent, it is possible, for example, to have a subtree as part of several trees.




Figure 6.3 *Highly coupled tree structure with parent pointers*

If many messages from one object to another require the same parameter, it may be better to permanently attach the parameter to the called object. Parameters are thin threads tying the object together, but, like Gulliver with the Lilliputians, enough thin threads can render an object incapable of being changed.

Figure 6.4 shows a single parameter.



Figure 6.4 *A single parameter introduces little coupling*

Illustrated with code:

```
Server s= new Server();
s.a(this);
```

Repeating the same parameter five times increases the coupling substantially:

```
Server s= new Server();
s.a(this);
s.b(this);
s.c(this);
s.d(this);
s.e(this);
```




Figure 6.5 Repeated parameter increases coupling

In this case, the two objects are better prepared to operate independently if the parameters are replaced by a pointer:

```
Server s= new Server(this);
s.a();
s.b();
s.c();
s.d();
s.e();
```




Figure 6.6 Reference reduces coupling

Collecting Parameter

Calculations that gather results from many method invocations need some way to merge those results. One way is to return a value from all of the methods. This works if the value is simple, like an integer.

```
Node
int size() {
    int result= 1;
    for (Node each: getChildren())
        result+= each.size();
    return result;
}
```

When merging the results is more complicated than simple addition, it is more direct to pass a parameter that will collect the results. For example, a collecting parameter helps when linearizing a tree:

```
Node
asList() {
    List results= new ArrayList();
    addTo(results);
    return results;
}
addTo(List elements) {
    elements.add(getValue());
    for (Node each: getChildren())
        each.addTo(elements);
}
```

Some other examples of more-complicated collecting parameters are the `GraphicsContext` that is passed around a tree of widgets or the `TestResult` that is passed around a tree of tests in JUnit.

Optional Parameter

Some methods can take a parameter or supply a default parameter if one isn't present. In such cases put the mandatory parameters first in the parameter list and add the optional parameters at the end. This makes as many of the parameters as possible the same, and the optional parameters appear as alternatives at the end.

The `ServerSocket` constructors demonstrate optional parameters. The basic constructor takes no arguments, but there is also a version with an optional port number and another version with an optional port number and an optional backlog length:

```
public ServerSocket()
public ServerSocket(int port)
public ServerSocket(int port, int backlog)
```

Languages with keyword parameters can express optional parameters more directly. Since Java has only positional parameters, whether a parameter is optional can only be expressed by convention. Some people call this the telescoping parameter pattern to provide a physical analogy for how the collection of parameters builds on each other.

Var Args

Some methods can take any number of a given type of parameter. A simple solution is to always pass a collection as the parameter. Callers of such methods, however, are cluttered by the presence of the intermediate collections:


```
Collection<String> keys= new ArrayList<String>();  
keys.add(key1);  
keys.add(key2);  
object.index(keys);
```

This problem is common enough that Java provides a mechanism to pass a variable number of arguments to a method. By declaring the method above as `method(Class... classes)`, the client can invoke the method with any number of arguments:

```
object.index(key1, key2);
```

Var args must be the last parameter. If a method has var args and optional arguments as described above, the optional arguments have to go before the var args.

Parameter Object

If a group of parameters is passed together to many methods, consider making an object whose fields are those parameters and passing the object instead. Once you have replaced the long parameter lists with the parameter object, see if there are bits of code using only the fields in the parameter object that you could turn into methods on the parameter object.

For example, it is common in Java graphics libraries to represent rectangles as independent `x`, `y`, `width`, and `height` parameters. Sometimes these four parameters are passed through many layers of method invocation, resulting in code that is longer and more difficult to read than it needs to be.

```
setOuterBounds(x, y, width, height);  
setInnerBounds(x + 2, y + 2, width - 4, height - 4);
```

Making the rectangle explicit as an object explains the code better:

```
setOuterBounds(bounds);  
setInnerBounds(bounds.expand(-2));
```

The introduction of the parameter object shortens the code, explains its intent, and provides a home for the algorithm for expanding and contracting a rectangle, which otherwise would have to be repeated everywhere it was required (with frequent errors when programmers forgot that the expansion factor has to be doubled for the width and height). Many powerful objects begin life as parameter objects.

While the primary motivation for introducing a parameter object is to improve readability, parameter objects can become important homes for logic. The fact that the data appears together in several parameter lists is a blatant

clue that they are strongly related. A class with its fixed list of fields is an explicit way of communicating “This set of data is strong related.”

The argument most often used against parameter objects is performance—allocating all those parameter objects takes time. In most cases this will not be a problem in practice. If object allocation becomes a bottleneck, the parameter object can be inlined, turned back into the explicit list of parameters, where necessary. The best code to optimize is readable, factored, and tested; parameter objects can contribute to these goals.

Constant

Sometimes you have data values that are needed several places within your program but which don’t change. If these values are known at compile time, store in variables declared `static final` and reference the variable throughout your program. It is common to give constants names consisting entirely of upper-case characters, to emphasize that they are not ordinary variables.

Part of the importance of using constants is that you avoid a whole class of errors when using them. If you have 5 embedded throughout your code and you decide to change 5 to 6, it is easy to miss an incidence. If 5 takes on two implicit meanings, say “draw a border” and “what follows is an acknowledgment packet”, then changing the constant is even more error-prone. The strongest reason to use constants, though, is because you can use the names of the constants to communicate what you mean by the value. Readers will be able to understand `Color.WHITE` much better than `0xFFFFFF`. If the encoding of color changes, code using the constant will not need to be changed.

A common use of constants is to communicate variations of a message in an interface. For example, to center text you could invoke `setJustification(Justification.CENTERED)`. One advantage of this style of API is that you can add new variants of existing methods by adding new constants without breaking implementors. However, these messages don’t communicate as well as having a separate method for each variation. In this style, the message above would be `justifyCentered()`. An interface where all invocations of a method have literal constants as arguments can be improved by giving it separate methods for each constant value.

Role-Suggesting Name

How do you decide what to call a variable? Many conflicting constraints come to bear on this question. I want to communicate my intent fully through my names, which often suggests long names. I’d like the names to be short to sim-

plify code formatting. Names will be read many times for each time they are typed, so the names should be optimized for readability, not ease of typing. Both the way the data in the variable is used and the role that data plays in the computation need to be expressed.

There are several pieces of information I need when I am trying to understand a variable. What is its purpose in the computation? How is the object referred to by the variable used? What is the scope and lifetime of the variable? How widely is the variable referenced?

Many variable naming schemes include type information in the names. Mine doesn't. What is the point of telling the compiler the types of variables over and over if I have to turn around and embed that same information again in the variable names? I can see including type information in variable names in languages that don't do much to prevent type errors, like C. Java provides ample support for avoiding type errors.

If I want to know the type of one of my variables, my IDE gives me quick feedback about it. Using short, composed methods also provides a quick reference to the most commonly used variables, locals and parameters.

Another facet of variables that readers need to understand is the scope of variables. Some variable naming practices encode the scope as a prefix to the name so `fCount` is a field and `lCount` is a local. Again, by composing relatively short methods I find that I am seldom confused by the scope of a variable. If I can't see a declaration of the variable here in this method, it is most likely a field (I use other techniques to avoid most static fields).

This leaves the role of the variable as the primary piece of information I try to communicate with my variable names, generally leading to short, clear names. If I have to struggle to find a name, it is generally because I don't understand the computation very well.

There are a few variable names that recur in my code:

- `result`—stores the object that will be returned from a function
- `each`—stores the individual elements of a collection while iterating (although I am becoming fond of using the singular form of the collection's name, for example `for (Node child: getChildren())`).
- `count`—stores counts

If I have multiple variables that I would like to give the same name, I qualify the name: `eachX` and `eachY` or `rowCount` and `columnCount`.

I am sometimes tempted to abbreviate words in variable names. This optimizes typing at the expense of reading. Since variables are read many times

for each time they are written, this is a false economy. Sometimes I am tempted to use several words for a variable, which makes the variable too long for comfortable typing. When this happens I look at the surrounding context. Why do I need so many words to distinguish this variable's role from the role of other variables? Often this leads me to simplify the design, allowing me to again write short variable names in good conscience.

To summarize, I communicate the role a variable plays through its name. Everything else important about the variable—its lifetime, scope, and type—can generally be communicated through context.

Declared Type

One of the features of Java and other pessimistically typed languages is the need to declare the types of your variables. Since you have to declare the type, you may as well use the declared type as an opportunity to communicate. Pick the declared type that communicates how the variable is to be used, not how it is implemented.

`List<Person> members= new ArrayList<Person>()` tells me that `members` is to be used like a `List`. I expect to see operations invoked like `get()` and `set()`, since what sets `List` apart from `Collection` is indexed access to elements.

When I was first drafting this pattern, I wrote it dogmatically. Then I tried the rigid rule that all variables should be declared as generally as possible. What I found was that the extra effort to generalize all types was not worth it. Sometimes a variable would be a `List`. Then I would pass it to a method where only `Collection` protocol was used. The inconsistency between the declarations was a bigger problem for readers than the lack of precision in just declaring it as a `List` everywhere it was used. Now I would say it more gently. It is useful to declare variables and methods with a general type where possible. Losing a little precision and generality to maintain consistency is a reasonable trade-off.

The best thing about generalizing declared types is that it opens up options for changing the concrete classes during later modifications. If I declare a variable as an `ArrayList`, I can't easily change it later to a `HashSet` the way I could if I declared it as a `Collection`. In general, the further a decision propagates, the less flexibility you have for future change. To maintain flexibility, allow as little information as possible to spread as narrowly as possible. There is more information in the statement "`members` contains an `ArrayList`" than in the statement "`members` contains a `Collection`".

Focusing on communication is a good heuristic for maintaining flexibility. Declared types are an example of this. When I say that a variable holds a `Collection`, I am speaking precisely. Communicating well provides the best flexibility.

Initialization

Before you can program, you need to know what you can count on. Being able to make accurate assumptions helps you focus on learning what you need to know. One issue about which it is helpful to be able to make assumptions is the state of variables. Initialization is the process of putting variables into a known state before they are used.

There are several issues in initializing variables. One is the desire to make initialization as declarative as possible. If initialization and declaration go together, there is one place to go for answers to questions about the variable. Another issue is performance. Variables that are expensive to initialize may need to be initialized some time after they come into existence. For example, in Eclipse, classes are loaded as late as possible to keep start-up time low.

Listed below are two initialization patterns: eager and lazy.

Eager Initialization

One style of initialization is to initialize a variable as soon as it comes into existence—when it is declared or when the object in which it lives is created (declaration or constructor). One advantage of eager initialization is that you can be assured that the variables are initialized before they are used.

Initialize variables in the declaration if possible. This puts the declared and actual types close together for readers.

```
class Library {  
    List<Person> members= new ArrayList<Person>();  
    ...  
}
```

Initialize fields in the constructor if they can't be initialized in the declaration:

```
class Point {  
    int x, y;  
    Point(int x, int y) {  
        this.x= x;  
        this.y= y;  
    }  
}
```

There is a certain symmetry to initializing all the fields of an object in the same place, either the declarations or the constructor. However, mixing the two styles doesn't seem to cause any confusion as long as the objects are kept to a reasonable size.

Lazy Initialization

Eager initialization works well when you don't mind paying the cost of computing a variable's value when the variable comes into existence. When the computation is expensive and you would like to defer the cost (perhaps because the variable may never be used), create a getter method and initialize the field when the getter is first called:

```
Library.Collection<Person> getMembers() {  
    if (members == null)  
        members= new ArrayList<Person>();  
    return members;  
}
```

Lazy initialization used to be a more common technique. Limitation in raw computing power was more often an issue. Lazy initialization is important when computational power is a limited resource. A resource-constrained environment like Eclipse, where start-up must be fast, uses lazy initialization to avoid loading plug-ins until they are about to be used.

A lazily initialized field is harder to read than one initialized eagerly. The reader has to look at least two places before understanding the implementation type of the field. When coding, you are storing information for future readers. Fortunately, there are only a few commonly asked questions, so a few techniques suffice to answer most of them. Lazy initialization says, "Performance is important here."

Conclusion

The state patterns talk about how to communicate decisions about representing the state in a program. The next chapter presents the other side of the coin, how to communicate decisions about the flow of control.

Chapter 7

Behavior

John Von Neumann contributed one of the primary metaphors of computing—a sequence of instructions that are executed one by one. This metaphor permeates most programming languages, Java included. The topic of this chapter is how to express the behavior of a program. The patterns are:

- Control Flow—Express computations as a sequence of steps.
- Main Flow—Clearly express the main flow of control.
- Message—Express control flow by sending a message.
- Choosing Message—Vary the implementors of a message to express choices.
- Double Dispatch—Vary the implementors of messages along two axes to express cascading choices.
- Decomposing Message—Break complicated calculations into cohesive chunks.
- Reversing Message—Make control flows symmetric by sending a sequence of messages to the same receiver.
- Inviting Message—Invite future variation by sending a message that can be implemented in different ways.
- Explaining Message—Send a message to explain the purpose of a clump of logic.
- Exceptional Flow—Express the unusual flows of control as clearly as possible without interfering with the expression of the main flow.
- Guard Clause—Express local exceptional flows by an early return.
- Exception—Express non-local exceptional flows with exceptions.

- **Checked Exception**—Ensure that exceptions are caught by declaring them explicitly.
- **Exception Propagation**—Propagate exceptions, transforming them as necessary so the information they contain is appropriate to the catcher.

Control Flow

Why do we have control flow in programs at all? There are languages like Prolog that don't have an explicit notion of a flow of control. Bits of logic float around in a soup, waiting for the right conditions before becoming active.

Java is a member of the family of languages in which the sequence of control is a fundamental organizing principle. Adjacent statements execute one after the other. Conditionals cause code to execute only in certain circumstances. Loops execute code repeatedly. Messages are sent to activate one of several subroutines. Exceptions cause control to jump up the stack.

All of these mechanisms add up to a rich medium for expressing computations. As an author/programmer, you decide whether to express the flow you have in mind as one main flow with exceptions, multiple alternative flows each of which is equally important, or some combination. You group bits of the control flow so they can be understood abstractly at first, for the casual reader, with greater detail available for those who need to understand them. Some groupings are routines in a class, some are by delegating control to another object.

Main Flow

Programmers generally have in mind a main flow of control for their programs. Processing starts here, ends there. There may be decisions and exceptions along the way, but the computation has a path to follow. Use your programming language to clearly express that flow.

Some programs, particularly those that are designed to work reliably in hostile circumstances, don't really have a visible main flow. These programs are in the minority, however. Using the expressive power of your programming language to clearly express little-executed, seldom-changed facts about your program obscures the more highly leveraged part of your program: the part that will be read, understood, and changed frequently. It's not that exceptional conditions are unimportant, just that focusing on expressing the main flow of the computation clearly is more valuable.

Therefore, clearly express the main flow of your program. Use exceptions and guard clauses to express unusual or error conditions.

Message

One of the primary means of expressing logic in Java is the message. Procedural languages use procedure calls as an information hiding mechanism:

```
compute() {  
    input();  
    process();  
    output();  
}
```

says, “For purposes of understanding this computation all you need to know is that it consists of these three steps, the details of which are not important at the moment.” One of the beauties of programming with objects is that the same procedure also expresses something richer. For every method, there is potentially a whole set of similarly structured computations whose details differ. And, as an extra added bonus, you don’t have to nail down the details of all those future variations when you write the invariant part.

Using messages as the fundamental control flow mechanism acknowledges that change is the base state of programs. Every message is a potential place where the receiver of the message can be changed without changing the sender. Rather than saying “There is something out there the details of which aren’t important,” the message-based version of the procedure says, “At this point in the story something interesting happens around the idea of input. The details may vary.” Using this flexibility wisely, making clear and direct expressions of logic where possible and deferring details appropriately, is an important skill if you want to write programs that communicate effectively.

Choosing Message

Sometimes I send a message to choose an implementation, much as a case statement is used in procedural languages. For example, if I am going to display a graphic in one of several ways, I will send a polymorphic message to communicate that a choice will take place at runtime.

```
public void displayShape(Shape subject, Brush brush) {  
    brush.display(subject);  
}
```

The message `display()` chooses the implementation based on the runtime type of the brush. Then I am free to implement a variety of brushes: `ScreenBrush`, `PostscriptBrush`, and so on.

Liberal use of choosing messages leads to code with few explicit conditionals. Each choosing message is an invitation to later extension. Each

explicit conditional is another point in your program that will require explicit modification in order to modify the behavior of the whole program.

Reading code that uses lots of choosing messages requires skill to learn. One of the costs of choosing messages is that a reader may have to look at several classes before understanding the details of a particular path through the computation. As a writer you can help the reader navigate by giving the methods intention-revealing names. Also, be aware of when a choosing message is overkill. If there is no possible variation in a computation, don't introduce a method just to provide the possibility of variation.

Double Dispatch

Choosing messages are good for expressing a single dimension of variability. In the example in “Choosing Message,” this dimension was the type of medium on which the shape was to be drawn. If you need to express two independent dimensions of variability, you can cascade two choosing messages.

For example, suppose I wanted to express that a Postscript oval was computed differently than a screen rectangle. First I would decide where I wanted the computations to live. The base computations seem like they belong in the Brush, so I will send a choosing message first to the Shape, then to the Brush:

```
displayShape(Shape subject, Brush brush) {
    shape.displayWith(brush);
}
```

Now each Shape has the opportunity to implement `displayWith()` differently. Rather than do any detailed work, however, they append their type onto the message and defer to the Brush:

```
Oval.displayWith(Brush brush) {
    brush.displayOval(this);
}
Rectangle.displayWith(Brush brush) {
    brush.displayRectangle(this);
}
```

Now the different kinds of brushes have the information they need to do their work:

```
PostscriptBrush.displayRectangle(Rectangle subject) {
    writer print(subject.left() + " " + "...+ " rect);
}
```

Double dispatch introduces some duplication with a corresponding loss of flexibility. The type names of the receivers of the first choosing message get

scattered over the methods in the receiver of the second choosing message. In this example, this means that to add a new `Shape`, I would have to add methods to all the `Brushes`. If one dimension is more likely to change than the other, make it the receiver of the second choosing message.

The computer scientist in me wants to generalize to triple, quadruple, quintuple dispatch. However, I've only ever attempted triple dispatch once and it didn't stay for long. I have always found clearer ways to express multi-dimensional logic.

Decomposing (Sequencing) Message

When you have a complicated algorithm composed of many steps, sometimes you can group related steps and send a message to invoke them. The intended purpose of the message isn't to provide a hook for specialization or anything sophisticated like that. It is just old-fashioned functional decomposition. The message is there simply to invoke the sub-sequence of steps in the routine.

Decomposing messages need to be descriptively named. Most readers should be able to gather what they need to know about the purpose of the sub-sequence from the name alone. Only those readers interested in implementation details should have to read the code invoked by the decomposing message.

Difficulty naming a decomposing message is a tip-off that this isn't the right pattern to use. Another tip-off is long parameter lists. If I see these symptoms, I inline the method invoked by the decomposing message and apply a different pattern, like `Method Object`, to help me communicate the structure of the program.

Reversing Message

Symmetry can improve the readability of code. Consider the following code:

```
void compute() {  
    input();  
    helper.process(this);  
    output();  
}
```

While this method is composed of three others, it lacks symmetry. The readability of the method is improved by introducing a helper method that reveals the latent symmetry. Now when reading `compute()`, I don't have to keep track of who is sent the messages—they all go to this.

```
void process(Helper helper) {  
    helper.process(this);  
}  
void compute() {  
    input();  
    process(helper);  
    output();  
}
```

Now the reader can understand how the `compute()` method is structured by reading a single class.

Sometimes the helper method invoked by a reversing message becomes important on its own. Sometimes, overuse of reversing messages can obscure the need to move functionality. If we had the following code:

```
void input(Helper helper) {  
    helper.input(this);  
}  
void output(Helper helper) {  
    helper.output(this);  
}
```

it would probably be better structured by moving the whole `compute()` method to the `Helper` class:

```
compute() {  
    new Helper(this).compute();  
}  
Helper.compute() {  
    input();  
    process();  
    output();  
}
```

Sometimes I feel silly introducing methods “just” to satisfy an “aesthetic” urge like symmetry. Aesthetics go deeper than that. Aesthetics engage more of your brain than strictly linear logical thought. Once you have cultivated your sense of the aesthetics of code, the aesthetic impressions you receive of your code is valuable feedback about the quality of the code. Those feelings that bubble up from below the surface of symbolic thought can be as valuable as your explicitly named and justified patterns.

Inviting Message

Sometimes as you are writing code, you expect that people will want to vary a part of the computation in a subclass. Send an appropriately named message to communicate the possibility of later refinement. The message invites programmers to refine the computation for their own purposes later.

If there is a default implementation of the logic, make it the implementation of the message. If not, declare the method abstract to make the invitation explicit.

Explaining Message

The distinction between intention and implementation has always been important in software development. It is what allows you to understand a computation first in essence and later, if necessary, in detail. You can use messages to make this distinction by sending a message named after the problem you are solving which in turn sends a message named after how the problem is to be solved.

The first example I saw of this was in Smalltalk. Transliterated, the method that caught my eye was this:

```
highlight(Rectangle area) {  
    reverse(area);  
}
```

I thought, “Why is this useful? Why not just call `reverse()` directly instead of calling the intermediate `highlight()` method?” After some thought, though, I realized that while `highlight()` didn’t have a computational purpose, it did serve to communicate an intention. Calling code could be written in terms of what problem they were trying to solve, namely highlighting an area of the screen.

Consider introducing an explaining message when you are tempted to comment a single line of code. When I see:

```
flags|= LOADED_BIT; // Set the loaded bit
```

I would rather read:

```
setLoadedFlag();
```

Even though the implementation of `setLoadedFlag()` is trivial. The one-line method is there to communicate.

```
void setLoadedFlag() {  
    flags|= LOADED_BIT;  
}
```

Sometimes the helper methods invoked by explaining messages become valuable points for further extension. It’s nice to get lucky when you can. However, my main purpose in invoking an explaining message is to communicate my intention more clearly.

Exceptional Flow

Just as programs have a main flow, they can also have one or more exceptional flows. These are paths of computation that are less important to communicate because they are less-frequently executed, less-frequently changed, or conceptually less important than the main flow. Express the main flow clearly, and these exceptional paths as clearly as possible without obscuring the main flow. Guard clauses and exceptions are two ways of expressing exceptional flows.

Programs are easiest to read if the statements execute one after another. Readers can use comfortable and familiar prose-reading skills to understand the intent of the program. Sometimes, though, there are multiple paths through a program. Expressing all paths equally would result in a bowl of worms, with flags set *here* and used *there* and return values with special meanings. Answering the basic question, “What statements are executed?” becomes an exercise in a combination of archaeology and logic. Pick the main flow. Express it clearly. Use exceptions to express other paths.

Guard Clause

While programs have a main flow, some situations require deviations from the main flow. The guard clause is a way to express simple and local exceptional situations with purely local consequences. Compare the following:

```
void initialize() {  
    if (!isInitialized()) {  
        ...  
    }  
}
```

with:

```
void initialize() {  
    if (isInitialized())  
        return;  
    ...  
}
```

When I read the first version, I make a note to look for an else clause while I am reading the then clause. I mentally put the condition on a stack. All of this is a distraction while I am reading the body of the then clause. The first two lines of the second version simply give me a fact to note: the receiver hasn’t been initialized.

If-then-else expresses alternative, equally important control flows. A guard clause is appropriate for expressing a different situation, one in which one of

the control flows is more important than the other. In the initialization example above, the important control flow is what happens when the object is initialized. Other than that, there is just a simple fact to notice, that even if an object is asked to initialize multiple times it will only execute the initialization code once.

Back in the old days of programming, a commandment was issued: each routine shall have a single entry and a single exit. This was to prevent the confusion possible when jumping into and out of many locations in the same routine. It made good sense when applied to FORTRAN or assembly language programs written with lots of global data where even understanding which statements were executed was hard work. In Java, with small methods and mostly local data, it is needlessly conservative. However, this bit of programming folklore, thoughtlessly obeyed, prevents the use of guard clauses.

Guard clauses are particularly useful when there are multiple conditions:

```
void compute() {
    Server server= getServer();
    if (server != null) {
        Client client= server.getClient();
        if (client != null) {
            Request current= client.getRequest();
            if (current != null)
                processRequest(current);
        }
    }
}
```

Nested conditionals breed defects. The guard clause version of the same code notes the prerequisites to processing a request without complex control structures:

```
void compute() {
    Server server= getServer();
    if (server == null)
        return;
    Client client= server.getClient();
    if (client == null)
        return;
    Request current= client.getRequest();
    if (current == null)
        return;
    processRequest(current);
}
```

A variant of guard clause is the `continue` statement used in a loop. It says, “Never mind this element. Go on to the next one.”

```
while (line = reader.readLine()) {  
    if (line.startsWith('#') || line.isEmpty())  
        continue;  
    // Normal processing here  
}
```

Again, the intent is to point out the (strictly local) difference between normal and exceptional processing.

Exception

Exceptions are useful for expressing jumps in program flow that span levels of function invocation. If you realize many levels up on the stack that a problem has occurred—a disk is full or a network connection has been lost—you may only be able to reasonably deal with that fact much lower down on the call stack. Throwing an exception at the point of discovery and catching at the point where it can be handled is much better than cluttering all the intervening code with explicit checks for all the possible exceptional conditions, none of which can be handled.

Exceptions cost. They are a form of design leakage. The fact that the called method throws an exception influences the design and implementation of all possible calling methods until the method is reached that catches the exception. They make it difficult to trace the flow of control, since adjacent statements can be in different methods, objects, or packages. Code that could be written with conditionals and messages, but is implemented with exceptions, is fiendishly difficult to read as you are forever trying to figure out what more is going on than a simple control structure. In short, express control flows with sequence, messages, iteration, and conditionals (in that order) wherever possible. Use exceptions when not doing so would confuse the simply communicated main flow.

Checked Exceptions

One of the dangers of exceptions is what happens if you throw an exception but no one catches it. The program terminates, that's what happens. But you'd like to control when the program terminates unexpectedly, printing out information necessary to diagnose the situation and telling the user what has happened.

Exceptions that are thrown but not caught are an even bigger risk when different people write the code that throws the exception and the code that catches the exception. Any missed communication results in an abrupt and impolite program termination.

To avoid this situation, Java has checked exceptions. These are declared explicitly by the programmer and checked by the compiler. Code that is subject to having a checked exception thrown at it must either catch the exception or pass it along.

Checked exceptions come with considerable costs. First is the cost of the declarations themselves. These can easily add 50% to the length of method declarations and add another thing to read and understand along the levels between the thrower and catcher. Checked exceptions also make changing code more difficult. Refactoring code with checked exceptions is more difficult and tedious than code without, even though modern IDEs reduce the burden.

Exception Propagation

Exceptions occur at different levels of abstraction. Catching and reporting a low-level exception can be confusing to someone who is not expecting it. When a web server shows me an error page with stack trace headed by a `NullPointerException`, I'm not sure what I'm supposed to do with the information. I'd rather see a message that said, "The programmer did not consider the scenario you have just presented." I wouldn't mind if the page also provided a pointer to further information that I could send to a programmer so he could diagnose the problem, but presenting me with untranslated details isn't helpful.

Low-level exceptions often contain valuable information for diagnosing a defect. Wrap the low-level exception in the higher-level exception so that when the exception is printed, on a log for example, enough information is written to help find the defect.

Conclusion

Control flows between methods of a program built from objects. The next chapter describes using methods to express the concepts in a computation.

This page intentionally left blank

Chapter 8

Methods

Logic is divided into methods, not smooshed together in one big lump. Why? What problems can be solved by introducing a new one of those pieces? What is the point of having methods at all? Conceptually, at least, you could organize any program as one gigantic routine with control jumping every which way. While this was how early programs were structured (with occasional recent reversions), the big-lump-of-logic suffers from problems. The most serious is that it is difficult to read. With one giant routine it is difficult to distinguish important parts from less important parts. It is difficult to understand part of the program now and leave some details for later. It is difficult to separate what is important for invokers of some functionality from what is important for those who need to modify that functionality. The second problem is that most problems encountered during programming are not unique. Rather than implement everything from scratch each time, it is handy (and productive) to be able to simply invoke a previous solution. The gigantic routine provides no handy way to refer to parts for later reuse.

Dividing the logic of a program into methods gives you a way to say, “These bits of logic aren’t closely connected.” Dividing the methods into classes and the classes into packages carries this communication further. Putting this code in this method and that code in that method signals readers that the two bits of code aren’t intimately related. They can be read and understood separately. Furthermore, the naming of methods gives you a chance to communicate to the reader what the purpose of this bit of the computation is, regardless of its implementation. Readers can often glean what they need from reading the names of methods alone.

Methods also neatly solve the reuse problem. When you are writing a new routine and you need a bit of logic that already exists as a method, you can invoke that method.

Dividing a large computation into methods is conceptually simple: put the pieces that go together together and the pieces that go apart apart. In practice,

though, you will spend time, energy, and creativity first figuring out what goes together and what doesn't, and second figuring out how best to make the division. What is a good division in the moment may not work when you change the logic of the system later on. The divisions need to be ones that simplify your overall workload. Knowing which division will work best comes through experience. Here are some of my hints from my experience.

The common issues in dividing a program into methods are the size, purpose, and naming of the methods. If you make too many too small methods, readers will have a hard time following your fragmented expression of ideas. Too few methods leads to duplication and the attendant loss of flexibility. There are many cliché tasks in programming, and creating a new method is a common step towards accomplishing many of these tasks. Methods that solve one of these recurring problems are generally easy to name. Naming methods that solve unique problems is harder but important for readers.

Here are the method-related patterns:

- **Composed Method**—Compose methods out of calls to other methods.
- **Intention-Revealing Name**—Name methods after what they are intended to do.
- **Method Visibility**—Make methods as private as possible.
- **Method Object**—Turn complex methods into their own objects.
- **Overridden Method**—Override methods to express specialization.
- **Overloaded Method**—Provide alternative interfaces to the same computation.
- **Method Return Type**—Declare the most general possible return type.
- **Method Comment**—Comment methods to communicate information not easily read from the code.
- **Helper Method**—Create small, private methods to express the main computation more succinctly.
- **Debug Print Method**—Use `toString()` to print useful debugging information.
- **Conversion**—Express the conversion of one type of object to another cleanly.
- **Conversion Method**—For simple, limited conversions, provide a method on the source object that returns the converted object.

- **Conversion Constructor**—For most conversions, provide a method on the converted object's class that takes the source object as a parameter.
- **Creation**—Express object creation clearly.
- **Complete Constructor**—Write constructors that return fully formed objects.
- **Factory Method**—Express more complex creation as a static method on a class rather than a constructor.
- **Internal Factory**—Encapsulate in a helper method object creation that may need explanation or later refinement.
- **Collection Accessor Method**—Provide methods that allow limited access to collections.
- **Boolean Setting Method**—If it helps communication, provide two methods to set boolean values, one for each state.
- **Query Method**—Return boolean values with methods named asXXX.
- **Equality Method**—Define `equals()` and `hashCode()` together.
- **Getting Method**—Occasionally provide access to fields with a method returning that field.
- **Setting Method**—Even less frequently provide the ability to set fields with a method.
- **Safe Copy**—Avoid aliasing errors by copying objects passed in or out of accessor methods.

Composed Method

Compose methods out of calls to other methods, each of which is at roughly the same level of abstraction.

One of the signs of a poorly composed method is a mixture of abstraction levels:

```
void compute() {  
    input();  
    flags |= 0x0080;  
    output();  
}
```

Code like this is jarring to the reader. Code is easier to understand when it flows, and abruptly shifting abstraction levels breaks flow. What is that bit

twiddling in there? I ask myself when reading the above method. What does it mean?

One objection to the use of lots of little methods is the performance penalty imposed by the invocation of all those methods. As I was writing this I wrote a little benchmark program that compared a million loop iterations with a million messages. The overhead was 20-30% on average, not enough to affect the performance of most programs. The combination of faster CPUs and the strongly localized nature of performance bottlenecks makes code performance an issue best left until you can gather statistics from realistic data sets.

How long should a method be? Some people recommend numerical limits, like less than a page or 5-15 lines. While it may be true that most readable code satisfies such a limit, limits beg the question “why?” Why do chunks of logic work out best when they are about so big?

Code readers need to solve several problems that provide opposing influences on the size of methods. When reading for overall structure, seeing lots of code at once is valuable. The white space in the method provides clues to the overall structure and complexity of the code. Are there conditionals and loops? How deeply nested are the control structures? How much work is necessary to accomplish the task implied by the name of the method?

The same big method that helped me orient myself becomes a hindrance when I turn to trying to understand the code in detail. I can only usefully hold one brainful of detail at a time, and a thousand-line method contains way more than one brainful. To understand details, I want closely related details gathered together and segregated from irrelevant details.

Simultaneously supporting browsing and digesting is the challenge of the code author who is dividing logic into methods. I find that my code reads best when I break it into relatively small methods (at least by C standards). The trick is recognizing when I have relatively independent sets of details that can be deferred to supporting methods. Sometimes I have too many details for easy comprehension but that are not easy to partition. In that case I create a method object to give all the details a place to be organized.

Another issue in choosing method size is specialization. Right-sized methods can be overridden whole, without having to copy code down to the subclass and edit it nor having to override two methods for one conceptual change.

Compose methods based on facts, not speculation. Get your code working, then decide how it should be structured. If you spend a lot of time structuring your code beforehand, you’ll just have to undo all that work and redo it when you learn something during implementation. When I have all the details of the logic laid out in front of me, it is much easier to compose methods sensibly. Sometimes I think I know how methods should be composed, but when I get

the logic divided up I find that the result is hard to read. In such cases I find it helpful to inline all the methods until I have one gigantic method again and re-partition it based on my recent experience.

Intention-Revealing Name

Methods should be named for the purpose a potential invoker might have in mind for using the method. There are other bits of information you might like to convey in the name—implementation strategy, for example. However, communicate intent in the name and communicate other information about the method in other ways.

Implementation strategy is the extraneous information most often included in method names. For example:

```
Customer.linearCustomerSearch(String id)
```

This might seem superior to:

```
Customer.find(String id)
```

because it communicates more about the method. However, the goal as a code author is not to simply blurt out everything you can about your program as quickly as you can. Sometimes, restraint is required. Unless the implementation strategy is relevant to users, leave it out of the name. The curious can look at the body of the method to see how it is implemented. Even if `Customer` offered both linear and hashed lookup, it would be better to communicate the distinction from the invoker's perspective through the names:

```
Customer.find(String id)  
Customer.fastFind(String id)
```

(Actually, in this case it would probably be better to just offer one `find()` that satisfied all users, but that's a different story.) Whether the fast version of `find()` is implemented with a hash table or a tree is not really relevant to users of the method.

Think about methods' names based on how they look in calling code. That's where readers are likely to first encounter the name. Why was this method invoked and not some other? That is a question that can profitably be answered by the name of the method. The calling method should be telling a story. Name methods so they help tell the story.

If you are implementing methods by analogy with an existing interface, give your methods the same names as those used in the interface. If you have a special kind of iterator, call your methods `hasNext()` and `next()`, even if you don't

formally implement the Iterator interface. If your methods are only kind of similar, first think about whether you are using the right metaphor, then express the differences as prefixes to the method name.

Method Visibility

The four levels of visibility—public, package, protected, private—each say something different about your intentions for a method.

The two big conflicting constraints in method visibility are the need to have some functionality revealed to outside users and the need to maintain future flexibility. The more methods that are revealed, the harder it is to change the interface to an object should you need to. In developing JUnit, Erich Gamma and I often disagree about the visibility of methods. My Smalltalk background suggests that making methods visible is potentially valuable to clients. Erich's Eclipse experience has taught him to value the flexibility that comes from revealing as little as possible. I'm slowly coming around to his point of view.

You have two costs to balance in choosing visibility. One is the cost of future flexibility. A very narrow interface makes future changes easier. The other cost is the cost to invoke your object. A too-narrow interface leaves all clients performing more work than necessary to use your object. Balancing these costs is central to making good visibility decisions.

My general strategy for visibility is to restrict it as much as possible. If that were all there was to picking visibility, it would be easy. A tool could assign method visibility. The real challenge comes when you are no longer working with certain knowledge, when code out of your direct control begins invoking your methods. Then you have to speculate, to decide which methods are going to be public or protected, committing you to maintaining them or paying a substantial cost to change them.

- **Public:** when you declare a method public you are saying that you believe it is useful outside of the package in which it is declared. Making a method public means you accept responsibility for maintaining it, either by leaving it unchanged, by fixing all callers if you change it, or by at least notifying the programmers who call it.

```
public Object next();
```

This declaration says that now and for the foreseeable future, `next()` will be available to clients.

- **Package:** package visibility is a statement that this method is useful to objects in this package but you aren't willing to commit to making it avail-

able to outside objects. This is a kind of odd statement—other objects need this method, but not all other objects, just mine. Treat package method visibility as a suggestion that either functionality should be moved around so the method can be made less visible or perhaps the method is more widely useful than you suspected and it may be worth the cost of making it public.

- **Protected:** protected visibility is only useful when offering code to be reused by subclassing. While it seems more restrictive than package visibility, the two are really orthogonal, since subclasses outside the package can see and invoke protected methods.
- **Private:** private methods are the ultimate in future flexibility since you are guaranteed that you will be able to find and change all callers, whether outsiders use and extend your code or not. By making a method private, you are saying that the value of this method to outsiders is not worth the cost of making it more widely available.

Slowly reveal methods, beginning with the most restrictive visibility that will work and revealing them as necessary. If a method no longer needs to be so visible, reduce its visibility. Reducing visibility only works when you have access to all the callers of your code, so you can be certain that you are not breaking a client by eliminating from view a method they relied on. I often notice that methods that I initially thought of as private become valued members of an interface once I begin using an object in new ways.

Declaring methods `final` is similar to choosing their visibility. Declaring a method `final` states that while you don't mind people using this method, you won't allow anyone to change it. If the invariants maintained by the method are complicated and subtle enough, this level of self-protection may be justified. You pay for the assurance that no one will accidentally break your object by eliminating the possibility that someone could profitably override it, and instead will have to do more work to get their job done some other way. I don't use `final` myself, and I have occasionally been frustrated when encountering `final` methods when I had a legitimate reason to override a method.

Declaring a method `static` makes it visible even if the caller does not have access to an instance of the class (subject to modification by the other visibility keywords). Static methods are limited in that they can't rely on any instance state, so they aren't a good repository for complex logic. Static methods can be inherited, but once overridden the superclass method cannot be invoked. One good use of static methods is as a replacement for constructors.

Method Object

This is one of my favorite patterns, probably because I use it so infrequently but the results are spectacular when I do. Creating method object can help you turn a tangled mass of code packed into an impossible method into readable, clear code that gradually reveals details to readers. This is a pattern I apply after I have some code working, and the more complex the method the better.

To create a method object, look for a long method with lots of parameters and temporary variables. Trying to extract any part of the method would result in long parameter lists in difficult-to-name sub-methods. Here are the steps to create a method object (since this refactoring is not yet supported automatically as I write this):

1. Create a class named after the method. For example, `complexCalculation()` becomes `ComplexCalculator`.
2. Create a field in the new class for each parameter, local variable, and field used in the method. Give these fields the same names as they have in the original method (you can fix names later).
3. Create a constructor that takes as parameters the method parameters of the original method and the fields of the original object used by the method.
4. Copy the method into a new method, `calculate()`, in the new class. The parameters, locals, and fields used in the old method become field references in the new object.
5. Replace the body of the original method with code that creates an instance of the new class and invokes `calculate()`. For example:

```
complexCalculation() {  
    new ComplexCalculator().calculate();  
}
```

6. If fields were set in the original method, set them after `calculate()` returns:

```
complexCalculation() {  
    ComplexCalculator calculator= new ComplexCalculator();  
    calculator.calculate();  
    mean= calculator.mean;  
    variance= calculator.variance;  
}
```

Make sure the refactored code works just like the old code did. Now the fun begins. The code in the new class is easy to refactor. You can extract methods and never have to pass any parameters because all the data used by the method

is stored in fields. Often, once you begin extracting methods you will discover that some variables can be demoted from fields to locals. Similarly, there may be information that can be passed into a single method as a parameter instead of being stored in a field. Once you begin extracting methods, you may find that common sub-expressions that were hard to isolate before become useful helper methods with meaningful names.

Sometimes, by the time I suspect that a method object may be called for, the original method has already been sliced up. In such a case, inline all the sub-methods so you have it all in one place before you begin. A clear indication that you need to do more inlining before making the method object is the need to call methods in the original object. Back up, inline them, and start over.

Overridden Method

One of the beauties of programming with objects is the variety of ways they provide to express the differences between similar calculations. Overridden methods are a clear way to express a variation. Methods declared abstract in a superclass are a clear invitation to specialize a calculation, but any method not declared final is a candidate for expressing a variation on an existing calculation. Well-composed methods in the superclass provide a multitude of potential hooks on which you can hang your own code. If the superclass code is in small, cohesive chunks, then you'll be able to override whole methods.

Overriding a method is not either/or. You can execute the subclass' code and the superclass' code by invoking `super.method()`;. Only do this to invoke the method of the same name. If your subclass explicitly chooses to sometimes invoke its own code and sometimes to invoke the superclass code for a variety of methods, the class will be hard to follow and easy to break accidentally. If you feel the need for invoking different superclass methods, you can improve the code by restructuring the control flow until bouncing back and forth between subclass and superclass is no longer necessary.

Superclass methods that are too large create a dilemma: copy code down to the subclass and edit it or find another way to express the variation? The problem with copying is that someone can come along later and change the superclass code you've copied, breaking your code without you (or them) knowing it.

Overloaded Method

When you declare the same method with different parameter types, you are saying, "Here are alternative formats for the parameters to this method." An

example is a method that can take a `String` representing a file name or an `OutputStream`, giving users who want to talk in terms of file names a simple interface but preserving flexibility for those users who want to pass in an already-formed stream (for testing, for example). Overloaded methods relieve the caller of the responsibility of converting parameters if there are several legitimate ways of passing the parameters.

A variant of overloading is using the same method name with different numbers of parameters. The problem with this style of overloading is that readers who want to ask, “What happens when I invoke this method?” need to read not only the name of the method but also the parameter list before they know enough to figure out what happens as a result of the method invocation. If the overloading is complicated, readers need to understand subtle overload resolution rules to be able to statically determine which method will be invoked for given types of arguments.

Overloaded methods should all serve the same purpose, with the variation only in the parameter types. Different return types for different overloaded methods make reading the code too difficult. Better to find a new name for the new intention. Give different computations different names.

Method Return Type

The return type of a method signals first whether the method is a procedure that works by side-effect or a function returning a particular type of object. The magic return type `void` allows Java to avoid a keyword to distinguish between procedures and functions.

Assuming you are writing a function, pick a return type to express your intention. Sometimes your intention is that the return type is specific, a concrete class or one of the primitive types. However, you’d like your methods to be as broadly applicable as possible, so pick the most abstract return type that expresses your intention. This preserves the flexibility for you to change the concrete return type should that become necessary in the future.

Generalizing the return type can also be a way for you to hide implementation details. For example, returning a `Collection` instead of a `List` can encourage users not to assume that elements are in a fixed order.

Return types are a common area for changes as you evolve your program. You may begin by returning a concrete class and then discover later that several related methods return different concrete classes, each of which do or should share a common interface. Expressing this similarity by declaring the common interface (if necessary) and returning the new interface from all the methods will help readers understand the similarity.

Method Comment

Express as much information as possible through the names and structure of the code. Add comments to express information that is not obvious from the code. Where they are expected, add javadoc comments to explain the purpose of methods and classes.

Many comments are completely redundant in code written to communicate. The cost of writing them and maintaining their consistency with the code is not worth the value they bring.

Method comments are at an awkward level of abstraction. If there are constraints between two methods (one must be called before the other, for example), where does the comment go? Comments must be kept up-to-date separately with the code, and there is no immediate feedback when comments are no longer valid.

Automated tests can communicate information that doesn't fit naturally in method comments. In the example above, I can write a test that ensures that the appropriate exception is thrown if the methods are invoked in the wrong order (though I would prefer to eliminate or encapsulate this constraint). Automated tests have many advantages. Writing them is a valuable design exercise, especially when done before implementation. If the tests run, they are consistent with the code. Automated refactoring tools can help keep tests updated at low cost.

All that said, communication is still the paramount value in these implementation patterns. If a method comment is the best possible medium for communication, write a good comment.

Helper Method

Helpers are one consequence of composed methods. If you are going to divide big methods into several smaller ones, you need those smaller methods. These are the helpers. Their purpose is to make larger-scale computations more readable by hiding temporarily irrelevant details and giving you a chance to express your intention through the name of the method. Helpers are typically declared private, moving to protected if the class is intended to be refined by subclassing.

You may create a method as a private helper only to discover that external users wish to invoke it. If the method is useful internally, there is a chance that it will be useful externally as well. Even if your little helper never “graduates”, however, it is still valuable as a point of communication.

Helpers tend to be short, but they can be too short. Just today I eliminated a helper that returned only a class' constructor. I find that:

```
return testClass.getConstructor().newInstance();
```

communicates as well as:

```
return getTestConstructor().newInstance();
```

However, that helper method might still be justified if subclasses were overriding how the constructor was computed.

Eliminate helpers (at least temporarily) when the logic of a method becomes unclear. Inline all the helper methods, take a fresh look at the logic, and re-extract methods that make sense.

A final purpose for helper methods is eliminating common sub-expressions. If you call a helper method every place in a class you need a certain little calculation, then changing that expression is easy. If the same one or two or three lines is duplicated throughout the object, not only have you lost the opportunity to communicate its purpose through a well-chosen method name, it is difficult to change.

Debug Print Method

There are many potential reasons to render an object as a string. You might want to present the object to a user, store the object for later retrieval, or present the internals of the object to a programmer.

The `Object` interface is fairly narrow, containing eleven methods. One of these methods, `toString()`, renders the receiver as a string, but to what purpose? One temptation is to satisfy several purposes at once. However, these compromises seldom work out. What a bond trader, a programmer, and a database want to know about an object are different.

There is leverage in investing in high-quality debug printing. To discover an important internal detail about an object can require half a minute of mouse clicks. Render that same detail in `toString()` and the same information is available in a second with a single click. I would rather be figuring out my program during debugging than navigating objects in the development environment. In intense debugging sessions, maintaining focus can save minutes or hours of effort.

Since `toString()` is public, it is subject to abuse. If an object doesn't support the needed protocol, people have been known to parse the print string to retrieve useful information. Such code is fragile because changing `toString()` is common. The best policy to prevent such abuse is to do your best to make sure your objects have all the protocol clients need.

Therefore, override `toString()` when you need to provide a programmer-friendly rendering of an object. Write other string renderings as other methods on the object or in separate classes.

Conversion

Sometimes you have object A and you need object B to pass to some further calculation. How do you represent this conversion from a source object to a destination object?

As with all these patterns, the goal of the conversion patterns is to communicate the programmer's intention clearly. However, there are some technical factors that influence which is the most effective way to express conversion. One of these is the number of conversions needed. If an object only needs to be converted to one other object, you can afford a simple approach. A potentially unbounded number of conversions requires a different approach. Another issue to consider is the dependencies between classes. It's not worth introducing a new dependency just to have a convenient expression of conversion.

The implementation of conversion is a whole separate issue. Sometimes you create a real object of the new type, copying information from the source object. Sometimes you can implement the interface of the destination object without copying information out of the source. As an alternative to conversion, sometimes you can just find a common interface for both objects and code to that.

Conversion Method

If you need to express conversion between objects of similar type, and there are a limited number of conversions, represent conversion as a method on the source object.

For example, suppose you want to implement cartesian and polar coordinates. To create a conversion method, you would implement:

```
class Polar {
    Cartesian asCartesian() {
        ...
    }
}
```

and vice versa. Notice that the return type of the conversion method is a class specific to the destination object. The point of conversion is to get an object with different protocol. Alternatively, you could implement `getX()` and `getY()` on

`Polar`, declare `Polar` and `Cartesian` as implementors of `Point`, and ignore conversion altogether.

Conversion methods have the advantage that they read nicely. They are quite popular (more than one hundred examples in *Eclipse*, for instance). However, to create one you need to be able to modify the source object's protocol to introduce one. They also introduce a dependency from the source object to the destination object. If such a dependency doesn't already exist, it's not worth introducing one just for the convenience of a conversion method. Finally, conversion methods become unwieldy when there are an unbounded number of potential conversions. A class with twenty different `asThis()`'s and `asThat()`'s is hard to read. Alternatively, you might change the clients so they can handle the source object instead of requiring conversion.

These disadvantages lead me to use conversion methods sparingly and only in situations where I am converting to objects of similar type. Otherwise, I use conversion constructors to express conversion.

Conversion Constructor

A conversion constructor takes the source object as a parameter and returns a destination object. A conversion constructor is useful when converting one source object into many destinations because the conversions don't all pile up in the source object.

For example, `File` supports a conversion constructor that converts a `String` representing the name of a file into an object suitable for reading, writing, and deleting. While it would be convenient to have `String.asFile()`, there is no end to the number of such conversions so it's better to have `File(String name)` and `URL(String spec)` and `StringReadStream(String contents)`. Otherwise, `String` would have an unbounded number of conversion methods.

If you need the freedom to implement conversion by returning something other than a concrete class, the conversion constructor can be expressed as a factory method returning a more general type (or placed on a different class than the one created by the method).

Creation

In the olden days (half a century ago), programs were big, undifferentiated masses of code and data. Control could flow from anywhere to anywhere. Data could be accessed from anywhere. Calculations, the original purpose of computers, occurred with (relatively speaking) lightning speed and perfect accuracy. Then people discovered an awkward fact: programs are written as much to be

changed as to be run. All this control jumping around and self-modifying code and data accessed from everywhere was great for execution, but it was terrible if you wanted to change the program later. And so began the long and halting road to find models of computation so a change *here* doesn't cause an unanticipated problem *there*.

Smaller programs are generally easier to modify than larger programs. One early strategy to make programs easier to change was to divide the big computer running a big program up into a bunch of smaller computers (objects) running little programs. Objects serve future change by providing an event horizon beyond which changes to the program have low cost.

This subdivision is for human purposes; to accommodate our fallible, changing, inventive minds, not for the good of the computer. The computer runs the same whether the code is one big ugly lump or a lovingly crafted network of mutually supportive objects. For human readers, creating an object makes a statement: some state goes together to support some computation, the details of which are irrelevant for now.

Using object creation expressively requires a balance between the need for clear and direct expression and the need for flexibility. The implementation patterns around creation provide techniques to express variations on the theme of “make me an object”.

Complete Constructor

Objects need certain information before they can compute. Communicate the prerequisites to potential users by providing constructors that return objects ready to compute. If there are multiple ways to set up an object, provide multiple constructors, each of which returns a well-formed object.

```
new Rectangle(0, 0, 50, 200);
```

Flexibility is sometimes better served by creating an object with a zero-argument constructor followed by a series of setting methods. However, this approach doesn't communicate what combinations of parameters are required for the correct operation of the object.

```
Rectangle box= new Rectangle();  
box.setLeft(0);  
box.setWidth(50);  
box.setHeight(200);  
box.setTop(0);
```

Can I get by without some of these parameters? There is no way to tell by reading the interface. If I see a four-argument constructor, though, I know all four arguments are required.

Constructors commit clients to a concrete class. At the time you write the code invoking the constructor, you may be satisfied that you want to use the concrete class. If you want to make your code more abstract, introduce a factory method. Even if you have a factory method, provide a complete constructor beneath it so curious readers can quickly understand what parameters are needed to create an object.

When implementing a complete constructor, funnel all the constructors to a single master constructor that does all the initialization. This ensures that all the variant constructors will create objects that satisfy all invariants required for proper operation and communicates those invariants to future modifiers of the class.

Factory Method

An alternative way to represent object creation is as a static method on the class. These methods have a couple of advantages over constructors: they can return a more abstract type (a subclass or an implementation of an interface) and they can be named after their intention, not just the class. However, factory methods add complexity, so they should be used when their advantages are valuable, not just as a matter of course.

Cast as a factory method, the `Rectangle` example would look like this:

```
Rectangle.create(0, 0, 50, 200);
```

If you intend something more complex than creating an object, like recording objects in a cache or creating a subclass that will be decided at runtime, then the factory method is helpful. However, as a reader I am always curious when I see a factory method. What else is going on there beyond object creation? I don't want to waste my readers' time, so if all that's happening is vanilla object creation, I express it as a constructor. If something else is happening at the same time, I introduce a factory method to tip curious readers to this fact.

A variant of a factory method is to gather related factory methods together as instance methods of a special factory object. This is useful when you have several concrete classes that vary at the same time. For example, each operating system could have a different factory object for creating the objects that make operating system calls.

Internal Factory

What do you do when the creation of a helper object is private but complex or subject to change by subclasses? Make a method which creates and returns the new object.

Internal factories are common in lazy initialization. The point of the getter method is stating that the variable is being initialized lazily:

```
getX() {  
    if (x == null)  
        x= ...;  
    return x;  
}
```

This is plenty for one method to communicate. If the calculation of *x* is at all complicated, it can profitably be deferred to an internal factory:

```
getX() {  
    if (x == null)  
        x= computeX();  
    return x;  
}
```

Internal factories are also an invitation to refinement by subclasses. A computation that uses the same algorithms on different data structures can be expressed through internal factories. Alternatively, you could pass the data structures as parameters to a helper object.

Collection Accessor Method

Suppose you have an object that contains a collection. How do you provide access to that collection? The simplest solution is providing a getting method for the collection:

```
List<Book> getBooks() {  
    return books;  
}
```

This gives clients maximum flexibility but creates a variety of problems. Internal state that depends on the contents of the collection can be invalidated behind your back if you return the whole collection. Providing such all-purpose access also passes up an opportunity to create a rich, meaningful protocol for your objects.

One alternative is to wrap the collection in an unmodifiable collection before returning it. Unfortunately, the wrapper only pretends to be a collection for the compiler. An exception is thrown if anyone tries to modify the wrapped

collection. Debugging such an error, especially in code in production, is expensive.

```
List<Book> getBooks() {  
    return Collections.unmodifiableList(books);  
}
```

Instead, offer methods that provide limited, meaningful access to the information in the collection.

```
void addBook(Book arrival) {  
    books.add(arrival);  
}  
int bookCount() {  
    return books.size();  
}
```

If clients need to iterate over the elements of the collection, provide a method that returns an iterator:

```
Iterator getBooks() {  
    return books.iterator();  
}
```

This prevents clients from modifying the collection, except for that pesky `remove()` operation on `Iterator`. If you want to ensure that clients don't change the contents of the collection, return an iterator that throws an exception if an element is removed. Again, only being notified of the error at runtime is risky and potentially expensive to debug.

```
Iterator<Book> getBooks() {  
    final Iterator<Book> reader= books.iterator();  
    return new Iterator<Book>() {  
  
        public boolean hasNext() {  
            return reader.hasNext();  
        }  
  
        public Book next() {  
            return reader.next();  
        }  
  
        public void remove() {  
            throw new UnsupportedOperationException();  
        }  
  
    };  
}
```

If you find yourself duplicating most of the collection protocols, it's likely you have a design problem. If your object did more work for its clients, it wouldn't have to offer so much access to its innards.

Boolean Setting Method

How best to offer protocol that sets boolean state? The simplest solution is a bare setting method:

```
void setValid(boolean newState) {  
    ...  
}
```

If clients need the flexibility of this access, this style is fine. However, when all the calls to the setting method are the constants `true` or `false`, you can offer a more expressive interface by providing two methods, one for each boolean value:

```
void valid() {...  
void invalid() {...
```

Code using this interface reads better, and in such code it is easier to statically discover where state is being set. However, if you see code like this:

```
...  
if (...boolean expression...)  
    cache.valid();  
else  
    cache.invalid();
```

go ahead and provide `setValidity(boolean)` instead.

Query Method

Sometimes an object needs to make a decision based on the state of another object. This isn't ideal, as the other object should generally make decisions for itself. However, when an object needs to offer decision criteria as part of its protocol, do so with a method whose name is prefixed with a form of "be" (like "is" or "was") or "have".

If one object has lots of logic that depends on the state of another object, it's a clue that the logic is misplaced. For example, if a method reads like this:

```
if (widget.isVisible())  
    widget.doSomething();  
else  
    widget.doSomethingElse();
```

then the widget is likely missing a method.

Try moving the logic and seeing if it reads more clearly. Sometimes these moves violate your preconceptions about which object is responsible for what part of a computation. Believing and acting on the evidence of your eyes generally improves the design. The result reads better and is more generally useful than a rigidly maintained picture drawn in advance of experience.

Equality Method

When two objects need to be compared for equality, for example because they are used as keys in a hash table, but their identity doesn't matter, implement `equals()` and `hashCode()`. Because two objects that are equal must have the same hash value, compute the hash only using data that is used in computing equality.

For example, if you are writing financial software you might have financial instruments with serial numbers. This could lead you to an equality method like this:

```
Instrument
public boolean equals(Object other) {
    if (! other instanceof Instrument)
        return false;
    Instrument instrument= (Instrument) other;
    return getSerialNumber().equals(instrument.getSerialNumber());
}
```

Note the guard clause at the beginning of the method. In theory, any two objects can be compared for equality, so your code should be prepared for this eventuality. If you know that cross-class comparisons are an indication of a programming error, eliminate the guard clause and allow a `ClassCastException` to be raised. Either that or throw an `IllegalArgumentException` inside the guard clause.

Since the serial number is the only information used in comparing equality, it is the only data that should be used in computing the hash value:

```
Instrument
public int hashCode() {
    return getSerialNumber().hashCode();
}
```

Note that for small data sets, 0 works just fine as a hash code.

The whole question of equality seemed more important twenty years ago. I can remember spending considerable time designing elaborate equality schemes. A cartoon that circulated at that time showed two diners sitting at a lunch counter. The first one says to the server, "I'll have what he's having,"

whereupon the server grabs the second diner's plate and plunks it down in front of the first diner.

Today's `equals()` and `hashCode()` are the vestigial remains of that concern with equality. If you are going to use them you need to follow the rules. Don't, and you'll end up with strange defects like putting an object into a collection and not being able to retrieve it immediately afterwards.

Another alternative to fussing with equality is to ensure that if two immutable objects are equal they are the same object. For example, allocating `Instruments` in a factory method enables this:

```
Instrument
static Instrument create(String serialNumber) {
    if (cache.containsKey(serialNumber))
        return cache.get(serialNumber);
    Instrument result= new Instrument(serialNumber);
    cache.put(serialNumber, result);
    return result;
}
```

Getting Method

One way of providing access to an object's state is to provide a method that returns that state. By convention, in Java these methods are prefixed with "get". For example,

```
int getX() {
    return x;
}
```

This convention is a form of metadata. I briefly tried naming getting methods simply after the variable they returned, but soon reverted. If readers found "getX" easier to read than "x", whatever my personal opinion, it was better to write what was expected.

How to write getting methods is not nearly as important or interesting a question as whether to write them, or at least make them visible. Following the principle of putting logic and data together, the need for public- or package-visible getting methods is a clue that logic should be elsewhere. Rather than writing the getting method, try moving the logic that uses the data instead.

There are a couple of exceptions to my aversion to visible getting methods. One is when I have a set of algorithms located in their own objects. Algorithms need access to data and need a getting method to receive it. Another is when I want a public method and it just so happens to be implemented by returning the value of a field. Finally, getting methods that will be invoked by tools will often have to be public.

Internal getting methods (private or protected) are useful for implementing lazy initialization or caching. As with all additional abstractions, these refinements are best deferred until needed.

Setting Method

If you need a method to set the value of a field, name it after the field prefixed with “set”. For example:

```
void setX(int newX) {  
    x= newX;  
}
```

I’m even more reluctant to make setting methods visible than getting methods. Setting methods are named for implementation, not intention. If a useful bit of an interface is best implemented by setting a field, that’s fine, but the name of the method should be written from the client code’s perspective. It’s better to understand what problem a client is solving by setting a value and provide a method that addresses that problem directly.

Using a setting method as part of the interface lets the implementation leak out:

```
paragraph.setJustification(Paragraph.CENTERED);
```

Naming the interface after the purpose of the method helps the code speak:

```
paragraph.centered();
```

even if the implementation of `centered()` is a setting method:

```
Paragraph:centered() {  
    setJustification(CENTERED);  
}
```

Setting methods used internally (private or protected) can be valuable, for instance, for updating dependent information. Our paragraph, for example, might need to redisplay whenever the justification changes. This could be implemented in the setting method:

```
private void setJustification(...) {  
    ...  
    redisplay();  
}
```

This use of a setting method acts like a simple constraint engine, ensuring that if *this* data changes, *that* dependent data over there changes to match (in

this case the internals of the paragraph and the information displayed on the screen).

Setting methods make code brittle. One principle is to avoid action at a distance. If object A relies on the details of object B's internal representation, a change to B's code will also require a change to A's code, not because A has changed in any fundamental way, but just because the assumptions on which A is written have changed. Better to move the logic and the data together. Perhaps A should own the data or B should offer more meaningful protocol.

As with getting methods, if you have a tool that needs to invoke setting methods, mark them "Tool Use Only" and make them public. Provide a more communicative and modular interface for humans.

Safe Copy

When using either a getting or setting method, you have potential aliasing problems, where two objects each assume they have exclusive access to a third. Aliasing problems are a symptom of deeper design problems, such as a lack of clarity about which object is responsible for which data, but you can avoid some defects by making a copy of an object before returning it or storing it:

```
List<Book> getBooks() {  
    List<Book> result= new ArrayList<Book>();  
    result.addAll(books);  
    return result;  
}
```

In this case it is probably better to provide collection accessor methods instead. However, if you have to provide access to the whole collection, this is a safe way to do it.

Setting methods can also be written with safe copies:

```
void setBooks(List<Book> newBooks) {  
    books= new ArrayList<Book>();  
    books.addAll(newBooks);  
}
```

I remember reviewing one banking system where safe copies were overused. Each accessor method (getter or setter) had two versions, one "safe" and the other not. To eliminate aliasing defects, huge object structures were copied every time a safe copying method was invoked. The system was too slow, so clients tended to use the unsafe version, which resulted in a host of aliasing defects. The underlying design problem, that objects weren't offering enough meaningful protocol, was never addressed.

Safe copying is strictly a palliative, to be used to protect code from uncontrolled outside access. It should rarely be part of the core semantics of an implementation. Immutable objects and composed methods provide simpler, more communicative interfaces that are less prone to error.

Conclusion

This chapter has described the patterns for creating methods. This concludes the patterns related to the Java language. The following chapter describes patterns for using the collection classes.

Chapter 9

Collections

I must say that I didn't expect this chapter to amount to much. When I started writing it, I thought I would end up with an API document—types and operations. The basic idea is simple: a collection distinguishes between objects in the collection and those not in the collection. What more was there to say?

What I discovered is that collections are a far richer topic than I ever suspected, both in their structure and the possibilities they offer for communicating intent. The concept of collections blends several different metaphors. The metaphor you emphasize changes how you use collections. Each of the collection interfaces communicates a different variation on the theme of a sack of objects. Each of the implementations also communicates variations, mostly with regard to performance. The result is that mastering collections is a big part of learning to communicate well with code.

Collection-like behavior used to be implemented by providing links in the data structure itself: each page in a document would have links to the previous and next pages. More recently, the fashion has swung to using a separate object for the collection which relates elements. This allows the flexibility to put the same object in several different collections without modifying the object.

Collections are important because they are a way of expressing one of the most fundamental kinds of variation in programming, the variation of number. Variation in logic is expressed with conditionals or polymorphic messages. Variation in the cardinality of data is expressed by putting the data into a collection. The precise details of that collection reveal much about the intention of the original programmer to a reader.

There is an old (by computer terms) saying that the only interesting numbers are 0, 1 and many (this saying was not written by a numerical analyst). If the absence of a field expresses “zero” and the presence of a field expresses “one”, then a field holding a collection is a way of expressing “many”.

Collections hover in a strange world halfway between a programming language construct and a library. Collections are so universally useful, and their

use is so well understood, that it almost seems time to have a mainstream language that allows statements like `plural unique Book books`; instead of the current `Collection<Book> books= new HashSet<Book>()`; Until collections are first-class language elements, it is important to know how to use the current collection library to express common ideas in straightforward ways.

The remainder of the chapter is divided into six parts: the metaphors behind collections, the issues to be expressed through the use of collections, the collection interfaces and what they mean to the reader, the collection implementations and what they say, an overview of functions available in the `Collections` class, and finally a discussion of extending collections through inheritance.

Metaphors

As suggested above, collections blend different metaphors. The first is that of a multi-valued variable. There is a sense in which a variable that refers to a collection is really a variable referring to several objects at the same time. Looked at this way, the collection disappears as a separate object. The collection's identity is not interesting, only the objects to which it refers. As with all variables, you can assign to a multi-valued variable (add and remove elements), retrieve its value, and send the variable messages (with the `for` loop).

The multi-valued variable metaphor breaks down in Java because collections are separate objects with identity. The second metaphor mixed into collections is that of objects—a collection is an object. You can retrieve a collection, pass it around, test it for equality, and send it messages. Collections can be shared between objects, although this creates the possibility of aliasing problems. Because collections are a set of related interfaces and implementations, they are open to extension, both with expanded interfaces and new implementations. So, just as collections “are” multi-valued variables, they also “are” objects.

The combination of the two metaphors makes for some strange effects. Because a collection is implemented as an object that can be passed around, you get the equivalent of call-by-reference, where instead of passing a variable's contents to a routine, you pass the variable itself. Changes to the variable's value are reflected in the calling routine. Call-by-reference went out of fashion in language design a couple of decades ago because of the possibility for unintended consequences. It was hard to debug programs when you couldn't be certain of all the places where a variable could be modified. Some of the conventions for programming with collections exist to avoid situations where it is hard to read the code and predict where a collection could be modified.

A third metaphor useful for thinking about collections is that of mathematical sets. A collection is a sack of objects just like a mathematical set is a sack of elements. A set divides the world into things in the set and things not in the set. A collection divides the world of objects into objects that are in the collection and objects that are not. Two basic operations on mathematical sets are finding their cardinality (the `size()` method of collections) and testing for inclusion (represented by the `contains()` method).

The mathematical metaphor is only approximate for collections. The other basic operations on sets—union, intersection, difference, and symmetric difference—are not directly represented by collections. Whether this is because these operations are intrinsically less useful or because they aren't used because they aren't available makes for an interesting debate.

Issues

Collections are used to express several orthogonal concepts in programs. In principle, you should express yourself as precisely as possible. With collections, this means using the most general possible interface as a declaration and the most specific implementation class. However, this is not an absolute rule. I carefully went through JUnit and generalized all the variable declarations. The result was a mess, because there was no uniformity. The confusion of having the same object declared as an `Iterable` in one spot, a `Collection` in another, and a `List` elsewhere made reading the code more difficult without much payoff. It was clearer to just declare every variable as a `List`.

The first concept expressed by collections is their size. Arrays (which are primitive collections) have a fixed size, set when the array is created. Most collections can change size after they are created.

A second concept expressed through collections is whether or not the order of elements is important. Calculations in which the elements affect each other or where external users of the calculation attach importance to order call for collections that preserve order. The order may be the order in which elements were added or it may be provided by some outside influence like lexicographic comparison.

Another issue to be expressed by collections is the uniqueness of elements. There are computations where the presence or absence of an element is sufficient, others where an element needs to be able to be present multiple times in a collection for the computation to be correct.

How are the elements accessed? Sometimes it is enough to iterate over the elements, doing some calculation with them one at a time. At other times it is important to be able to store and retrieve elements with a key.

Finally, performance considerations are communicated through choice of collection. If a linear search is fast enough, a generic `Collection` is good enough. If the collection grows too large it will be important to be able to test for or access elements by a key, suggesting a `Set` or `Map`. Time and space can both be optimized through the judicious selection of collections.

Sidebar: Performance

Most programmers don't have to worry about the performance of small-scale operations most of the time. This is a refreshing change from the old days, when performance tuning was daily business. However, computing resources are not infinite. When experience has shown that performance needs to be better and measurement has shown where the bottlenecks are, it is important to express performance-related decisions clearly. Many times, better performance results in less of some other quality in the code, like readability or flexibility. It is important to pay as little as possible for the needed performance.

Coding for performance can violate the principle of local consequences. A small change to one part of a program can degrade performance in another part. If a method works efficiently only if the collection it is passed can test for membership quickly, then an innocent substitution of `ArrayList` for `HashSet` elsewhere in the program can make the method intolerably slow. Distant consequences are another argument for coding carefully when coding for performance.

Performance is connected with collections because most collections can grow without limit. The data structure holding the characters I am typing right now needs to be able to hold millions of characters. I would like inserting the millionth character to be just as fast as inserting the first.

My overall strategy for performance coding with collections is to use the simplest possible implementation at first and pick a more specialized collection class when it becomes necessary. When I make performance-related decisions I try to localize them as much as possible even if that requires some changes to the design. Then, when the performance is good enough again, I stop tuning.

Interfaces

Readers of collection-based code are looking for answers to different questions when they look at the interfaces you have declared for your variables and the implementations you chose for those variables. The interface declaration tells the reader about the collection: whether the collection is in a particular order, whether there are duplicate elements, and whether there is any way to look up elements by key or only through iteration.

The interfaces described below are:

- **Array**—Arrays are the simplest and least flexible collection: fixed size, simple accessing syntax, and fast.
- **Iterable**—The basic collection interface, allowing a collection to be used for iteration but nothing else.
- **Collection**—Offers adding, removing, and testing for elements.
- **List**—A collection whose elements are ordered and can be accessed by their location in the collection (i.e., “give me the third element”).
- **Set**—A collection with no duplicates.
- **SortedSet**—An ordered collection with no duplicates.
- **Map**—A collection whose elements are stored and retrieved by key.

Array

Arrays are the simplest interface for collections. Unfortunately, they don’t have the same protocol as other collections, so it’s harder to change from an array to a collection than from one kind of collection to another. Unlike most collections, the size of an array is fixed when it is created. Arrays are also different as they are built into the language, not provided by a library.

Arrays are more efficient in time and space than other collections for simple operations. The timing tests I did to accompany writing this suggest that array access (i.e. `elements[i]`) is more than ten times faster than the equivalent `ArrayList` operation (`elements.get(i)`). (As these numbers vary substantially in different operating environments, if you care about the performance difference you should time the operations yourself.) The flexibility of the other collection classes makes them more valuable in most cases, but arrays are a handy trick to

be able to pull out when you need more performance in a small part of an application.

Iterable

Declaring a variable `Iterable` only says that it contains multiple values. `Iterable` is the basis for the loop construct in Java 5. Any object declared as `Iterable` can be used in a `for` loop. This is implemented by quietly calling the method `iterator()`.

One of the issues to be communicated when using collections is whether clients are expected to modify them. Unfortunately, `Iterable` and its helper, `Iterator`, provide no way to state declaratively that a collection shouldn't be modified. Once you have an `Iterator`, you can invoke its `remove()` method, which deletes an element from the underlying `Iterable`. While your `Iterables` are safe from having elements added, they can have elements removed without the object that owns the collection being notified.

As described in “Collection Accessor Method” on page 91, there are a few ways to ensure that a collection is not modified: wrapping it in a `unmodifiable` collection, creating a custom iterator that throws an exception when a client tries to modify the collection, or returning a safe copy.

`Iterable` is simple. It doesn't even allow you to measure the size of instances; all you can do is iterate over the elements. Sub-interfaces of `Iterable` provide more useful behavior.

Collection

`Collection` inherits from `Iterable`, but it adds methods to `add`, `remove`, `search` for and `count` elements. Declaring a variable or method as a `Collection` leaves many options for an implementation class. By leaving the declaration as vaguely specified as possible, you retain the freedom to change implementation classes later without having the change ripple through the code.

Collections are a bit like the mathematical notion of sets, except that the operations performing the equivalent of union, intersection, and difference (`addAll()`, `retainAll()`, and `removeAll()`) modify the receiver instead of returning newly allocated collections.

List

To `Collection`, `List` adds the idea that elements are in a stable order. An element can be retrieved by providing its index to the collection. A stable sequence is important when the elements of a collection interact with each other. For example, a

queue of messages that should be processed in their arrival order should be stored in a list.

Set

A Set is a collection that contains no duplicates (elements that would report that they are `equal()` to each other). This corresponds closely to the mathematical notion of set, although the metaphor is thin because adding an element to a Set modifies the collection rather than returning a new collection including the added element.

A Set discards information that most collections keep—the number of times an element appears. This is not a problem in cases where the presence or absence of an element is interesting but the number of times the element appears is not. For example, if I want to know who all the authors of books are in a library, I don't care how many books each author wrote. I just want to know who they are. A Set is an appropriate way to implement such a query.

The elements in a Set are in no particular order. Just because you iterate through them in a certain order once does not mean that the elements will appear in the same order the next time. This lack of predictable order is not a limitation in cases where the elements don't interact with each other.

Sometimes you want to store duplicates in a collection but remove them for a particular operation. Create a temporary Set and pass it to the operation:

```
printAuthors(new HashSet<Author>(getAuthors()));
```

SortedSet

The ordering and uniqueness attributes of collections are not mutually exclusive. At times you'd like to keep a collection in order but eliminate duplicates. SortedSet stores ordered-but-unique elements.

Unlike the ordering of a List, which is related to the order in which elements were added or by explicit indexes passed to `add(int, Object)`, the ordering in a SortedSet is provided by a Comparator. In the absence of an explicit order, the “natural order” of the elements is used. For example, strings are sorted in lexicographical order.

To compute the authors contributing to a library, you could use a SortedSet:

```
public Collection<String> getAlphabeticalAuthors() {
    SortedSet<String> results= new TreeSet<String>();
    for (Book each: getBooks())
        results.add(each.getAuthor());
    return results;
}
```

This example uses the default sorting of strings. If a `Book` had its author represented by an object, the code above might look like this:

```
public Collection<String> getAlphabeticalAuthors() {
    Comparator<Author> sorter= new Comparator<Author>() {
        public int compare(Author o1, Author o2) {
            if (o1.getLastName().equals(o2.getLastName()))
                return o1.getFirstName().compareTo(o2.getFirstName());
            return o1.getLastName().compareTo(o2.getLastName());
        }
    };
    SortedSet<String> results= new TreeSet<String>(sorter);
    for (Book each: getBooks())
        results.add(each.getAuthor());
    return results;
}
```

Map

The final collection interface is `Map`, which is a hybrid of the other interfaces. A `Map` stores values by key, but unlike a `List`, the key can be any object and not just an integer. The keys of a `Map` must be unique, a bit like sets, although the values can contain duplicates. The elements of a `Map` are in no particular order, also like a `Set`.

Because `Map` is not completely like any of the other collection interfaces, it stands alone, not inheriting from any of them. Maps are two collections at the same time; a collection of keys connected to a collection of values. You can't simply ask a `Map` for its iterator, because it is not clear whether you want an iterator over the keys, over the values, or over the pairs of keys-and-values.

Maps are useful for implementing two of the implementation patterns: extrinsic state and variable state. Extrinsic state suggests storing special-purpose data related to an object separately from the object itself. One way to implement extrinsic state is with a `Map` whose keys are the objects and whose values are the related data. In variable state, different instances of the same class store different data fields. To implement this, have the object hold a map which maps from strings (representing the names of the virtual fields) to the values.

Implementations

Choosing implementation classes for collections is primarily a matter of performance. As with all performance issues, it is best to pick a simple implementation to begin with and then tune based on experience.




Figure 9.1 Collection interfaces and classes

In this section, each interface introduces alternative implementations. Because performance considerations dominate the choice of implementation class, each set of alternatives is accompanied by performance measurements for important operations. Appendix, “Performance Measurement,” provides the source code for the tool I used to gather this data.

By far the majority of collections are implemented by `ArrayList`, with `HashSet` a distant second (~3400 references to `ArrayList` in Eclipse+JDK versus ~800 references to `HashSet`). The quick-and-dirty solution is to choose whichever of these classes suits your needs. However, for those times when experience shows that performance matters, the remainder of this section presents the details of the alternative implementations.

A final factor in choosing a collection implementation class is the size of the collections involved. The data presented below shows the performance of collections sized one to one hundred thousand. If your collections only contain one or two elements, your choice of implementation class may be different than if you expect them to scale to millions of elements. In any case, the gains available from switching implementation classes are often limited, and you'll need to look for larger-scale algorithmic changes if you want to further improve performance.

Collection

The default class to use when implementing a `Collection` is `ArrayList`. The potential performance problem with `ArrayList` is that `contains(Object)` and other operations that rely on it like `remove(Object)` take time proportional to the size of the collection. If a performance profile shows one of these methods to be a bottleneck, consider replacing your `ArrayList` with a `HashSet`. Before doing so, make sure that your algorithm is insensitive to discarding duplicate elements. When you have data that is already guaranteed to contain no duplicates, the switch won't make a difference. Figure 9.2 compares the performance of `ArrayList` and `HashSet`. (See Appendix A for the details of how I collected this information.)

List

To the `Collection` protocol, `List` adds the idea that the elements are in a stable order. The two implementations of `List` in common use are `ArrayList` and `LinkedList`.

The performance profiles of these two implementations are mirror images. `ArrayList` is fast at accessing elements and slow at adding and removing elements, while `LinkedList` is slow at accessing elements and fast at adding and removing elements (see Figure 9.3). If you see a profile dominated by calls to `add()` or `remove()`, consider switching an `ArrayList` to a `LinkedList`.

Set

There are three main implementations of `Set`: `HashSet`, `LinkedHashSet`, and `TreeSet` (which actually implements `SortedSet`). `HashSet` is the fastest but its elements are in no guaranteed order. A `LinkedHashSet` maintains elements in the order in which they were added, but at the cost of an extra 30% time penalty for adding and removing elements (see Figure 9.3). `TreeSet` keeps its elements sorted according to a `Comparator` but at the cost of making adding and removing elements or testing for an element take time proportional to $\log n$, where n is the size of the collection.




Figure 9.2 Comparing ArrayList and HashSet as implementations of Collection

Choose a `LinkedHashSet` if you need the order of elements to be stable. External users, for example, may appreciate getting elements in the same order each time.

Map

The implementations of `Map` follow a similar pattern to the implementations of `Set`. `HashMap` is the fastest and simplest. `LinkedHashMap` preserves the order of elements, iterating over the elements in the order in which they were inserted. `TreeMap` (actually an implementation of `SortedMap`) iterates over entries based on the order of the keys, but at the cost of making insertion and inclusion testing take time propor-




Figure 9.3 Comparing ArrayList and LinkedList

tional to $\log n$. Figure 9.5 summarizes the way performance scales for these Map implementations.

Collections

The utility class `Collections` is a library class that provides collection functionality that doesn't fit neatly into any of the collection interfaces. Here is a quick overview of what is available.




Figure 9.4 Comparing Set implementations.

Searching

The `indexOf()` operation takes time proportional to the size of the list. However, if the elements are sorted, binary search can find the index of an element in $\log_2 n$ time. Call `Collections.binarySearch(list, element)` to return the index of an element in the list. If the element does not appear in the list, a negative number will be returned. If the list is not sorted, the results are unpredictable.

Binary search only improves performance for lists with constant-time random access, like `ArrayList` (see Figure 9.6).




Figure 9.5 Comparing Map implementations

Sorting

Collections also provides operations to change the order of the elements of a list. `Reverse(list)` reverses the order of all the elements of the list. `Shuffle(list)` places the elements in random order. `Sort(list)` and `sort(list, comparator)` place the elements in ascending order. Unlike binary search, sorting performance is roughly the same for `ArrayList` and `LinkedList`, because the elements are first copied into an array, the array is sorted, and then the elements are copied back (run the timer test `Sorting` in Appendix A to verify this).




Figure 9.6 Comparing `indexOf()` and binary search

Unmodifiable Collections

As mentioned in the discussion of `Iterable` above, even the most basic collection interfaces allow collections to be modified. If you are passing a collection to untrusted code, you can ensure that it won't be modified by having `Collections` wrap it in an implementation that throws a runtime exception if clients try to modify it. There are variants that work with `Collection`, `List`, `Set`, and `Map`.

```
@Test(expected=UnsupportedOperationException.class)
public void unmodifiableCollectionsThrowExceptions() {
    List<String> l= new ArrayList<String>();
    l.add("a");
    Collection<String> unmodifiable= Collections.unmodifiableCollection(l);
    Iterator<String> all= unmodifiable.iterator();
    all.next();
    all.remove();
}
```

Single-Element Collections

If you have a single element and you need to pass it to an interface that expects a collection, you can quickly convert it by calling `Collections.singleton()`, which returns a `Set`. There are also variants that convert to a `List` or `Map`. All the collections returned are not modifiable.

```
@Test public void exampleOfSingletonCollections() {
    Set<String> longWay= new HashSet<String>();
    longWay.add("a");
    Set<String> shortWay= Collections.singleton("a");
    assertEquals(shortWay, longWay);
}
```

Empty Collections

Similarly, if you need to use an interface that expects a collection and you know you have no elements, `Collections` will create an unmodifiable empty collection for you.

```
@Test public void exampleOfEmptyCollection() {
    assertTrue(Collections.emptyList().isEmpty());
}
```

Extending Collections

I have often seen classes that extend one of the collection classes. A `Library` holding a list of books, for example, could be implemented by extending `ArrayList`:

```
class Library extends ArrayList {...}
```

This declaration provides implementations of `add()` and `remove()`, iteration, and the other collection operations.

There are several problems with extending collection classes to get collection-ish behavior. First, many of the operations offered by collections will be inappropriate for clients. For example, clients generally shouldn't be able to `clear()` a `Library` or `convert it toArray()`. At the very least, the metaphors are mixed and confusing. At worst, all these operations need to be disinherited by implementing them and throwing an `UnsupportedOperationException`. It's not a good trade-off to inherit a few useful lines of code but spend far more lines eliminating functionality you don't want. The second problem with inheriting from collection classes is that it wastes inheritance, a precious resource. To pick

up a few lines of useful implementation, you preclude using inheritance in some more highly leveraged way.

In such a situation, it is better to delegate to a collection rather than inherit from one:

```
class Library {  
    Collection<Book> books= new ArrayList<Book>();  
    ...  
}
```

With this design, you can reveal only those operations that make sense and you can give them meaningful names. You are free to use inheritance to share implementation with other model classes. If a `Library` offers access to books by several different keys, you can name the operations appropriately:

```
Book getBookByISBN(ISBN);  
Book getBookByID(UniqueID);
```

Only extend collections if you are implementing a general-purpose collection class, something that could be added to `java.util`. In all other cases, store elements in a subsidiary collection.

Conclusion

This chapter has described the patterns for using the collection classes. This concludes the patterns for Java and its collection classes. The preceding patterns were all written with a bias towards application development, where simplicity and ease of communication drive costs down but it is possible to change the design of the whole application at once. The following chapter describes how to modify these patterns when building frameworks, where complexity is acceptable if it preserves the ability to continue evolving the framework without being able to change all application code.