

Loop Parallelism with OpenMP

Riddle

Which assertion is right?

1. Processes have private heaps, and threads share a stack.
2. Processes share a heap, and threads have private stacks.
3. Processes have private heaps, and threads have private stacks.
4. Processes share a heap, and threads share a stack.
5. Processes are for lawyers, and you have no idea of this stack-heap thing.

From Program To Process (1)

```
int x;
int i, n;
n = 10;
x = 1;
for (i=0 ; i<n ; i++) {
    x = x * 2;
}
```

Source code

Assembly code

```
movl $10, -12(%rbp)
movl $1, -4(%rbp)
movl $0, -8(%rbp)
jmp L2
L3: sall -4(%rbp)
incl -8(%rbp)
L2: movl -8(%rbp), %eax
cmpl -12(%rbp), %eax
jl L3
movl -4(%rbp), %edx
.....
```

From Program To Process (2)

movl \$10, -12(%rbp)	movl \$10, -12(%rbp)
movl \$1, -4(%rbp)	movl \$1, -4(%rbp)
movl \$0, -8(%rbp)	movl \$0, -8(%rbp)
jmp L2	jmp L2
L3: sall -4(%rbp)	L3: sall -4(%rbp)
incl -8(%rbp)	incl -8(%rbp)
L2: movl -8(%rbp), %eax	L2: movl -8(%rbp), %eax
cmpl -12(%rbp), %eax	cmpl -12(%rbp), %eax
jl L3	jl L3
movl -4(%rbp), %edx	movl -4(%rbp), %edx
.....

From Program To Process (3)

```
OPCODE SRC DEST
4 +8 + 8
movl = 0010
jmp = 0100
jl = 0110
sall = 1000
incl = 0011
cmpl = 0101

direct value = 0 + value + 00
1 + 4 + 2
indirect = 1 + desloc + reg.
1 + 4 + 2

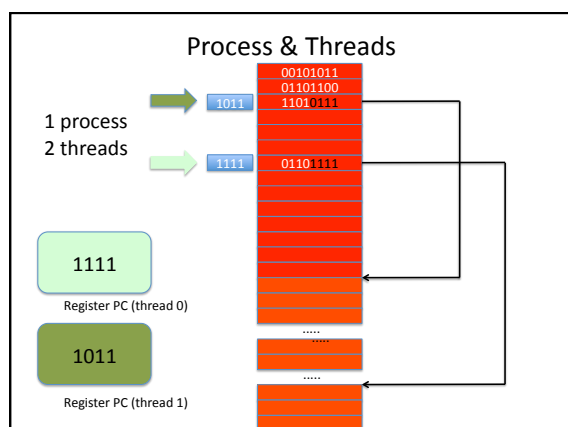
rbp = 00
eax = 01
edx = 10
```

0000	0010 0101000 1110000
0001	0010 0001000 1010000
0002	0010 0000000 1100000
0003	0100 0006
0004	1000 1010000
0005	0011 1100000
0006	0010 1100000 1000001
0007	0101 1110000 1000001
0008	0110 0004
0009	0010 1010000 1000010
0010

From Program To Process (4)

NOW this is an instant photography of the memory (RAM)

12	001100	0010 0101000 1110000
20	010100	0010 0001000 1010000
28	011010	0010 0000000 1100000
36	100100	0100 110100
44	101100	1000 1010000
52	110100	0011 1100000
60	111000	0010 1100000 1000001
68	0101 1110000 1000001
	0110 101100
	0010 1010000 1000010
	



Threads and segments

- Threads inside a process share the code segment.
 - A private PC points to a given instruction for each thread.
- Threads **share the heap** segment.
 - All dynamic vars are shared.
- Threads have **private stack** segments.
 - Vars local to a procedure are private.

Riddle

Which assertion is right?

1. Processes have private heaps, and threads share a stack.
2. Processes share a heap, and threads have private stacks.
3. **Processes have private heaps, and threads have private stacks.**
4. Processes share a heap, and threads share a stack.
5. Processes are for lawyers, and you have no idea of this stack-heap thing.

OpenMP

- A simple extension to sequential compilers to describe parallelism, based on:
 - Shared memory (threads),
 - Loops.
- Basic reference: <http://www.openmp.org>
- Designed by a **consortium** of companies
 - PGI, KAI, etc... (Compilers)
 - Intel, IBM, SGI, Sun, HP (HW + Compilers),
 - DoE, ASCI, NAG (users)
- You now get an OpenMP compiler directly “off the box”.

A pragma-based approach

```
double res[10000];

for (i=0 ; i<10000 ; i++)
  compute(&res[i]) ;
```

A pragma-based approach

```
double res[10000];

for (i=0 ; i<10000 ; i++)
  compute(&res[i]) ;
```

```
#include "openmp.h"
double res[10000];
#pragma omp parallel for
for (i=0 ; i<10000 ; i++)
  compute(&res[i]) ;
```

A simple flag at compile-time enables or disables the parallelism:

```
> gcc -fopenmp foo.c -o foo
> Export OMP_NUM_THREADS=4
> ./foo
```

Pragmas vs. Lib calls

In OpenMP, you find:

- Compile-time pragma
 - The compiler uses them to decide when to create threads, how to distribute the computations, etc.
 - Ex.: `#pragma omp num_threads(4)`
- Call to a runtime library
 - Executed at runtime!
 - Ex.: `omp_set_num_threads(4)`
 - Should be protected by `#ifdef _OPENMP ... #endif`

Basic OpenMP constructs

- `#pragma omp parallel`
 - Creates the threads.
 - Can be integrated (or not) with other pragmas
- `#pragma omp sections`
 - Declares sections of code to be run in different threads.

Example - sections

```
#pragma omp parallel
{ // now the threads are running
  #pragma omp sections
  foo1( );
  #pragma omp section
  bar2( );
  #pragma omp section
  foo2( );
}
```

Simple & efficient, yet...

- How do the threads deal with concurrent accesses in the memory?
- OpenMP enables the declaration of **private** vs. **shared** variables.
 - The pragma are complemented with `shared(x)`, `private(z)`, etc...
- Some variables are automatically private or shared:
 - Var. declared before the thread creation are shared.
 - Var allocated in the heap are shared (why)?
 - Local variables are private (why?)

One more example

```
double A[10000];
omp_set_num_threads(4);
#pragma omp parallel
{
  int th_id = omp_get_thread_num();
  compute(th_id, A);
}
printf("Done.");
```

One more example

```
double A[10000];
omp_set_num_threads(4);
#pragma omp parallel
{
  int th_id = omp_get_thread_num();
  compute(th_id, A);
}
printf("Done.");
```

`th_id` is private.

One more example

```
double A[10000];  
int th_id;  
omp_set_num_threads(4);  
#pragma omp parallel  
{  
    th_id = omp_get_thread_num();  
    compute(th_id, A);  
}  
printf("Done.");
```

th_id is shared.