

## Loop Parallelism with OpenMP

### OpenMP- A pragma-based approach

```
double res[10000];

for (i=0 ; i<10000 ; i++)
    compute(&res[i]) ;
```

```
#include "openmp.h"
double res[10000];
#pragma omp parallel for
for (i=0 ; i<10000 ; i++)
    compute(&res[i]) ;
```

A simple flag at compile-time enables or disables the parallelism:

- gcc -fopenmp foo.c -o foo
- Export OMP\_NUM\_THREADS=4
- ./foo

### One more example

```
double A[10000];
omp_set_num_threads(4);
#pragma omp parallel
{
    int th_id = omp_get_thread_num();
    compute(th_id, A);
}
printf("Done.");
```

### One more example

```
double A[10000];
omp_set_num_threads(4);
#pragma omp parallel
{
    int th_id = omp_get_thread_num();
    compute(th_id, A);
}
printf("Done.");
```

th\_id is private.

### One more example

```
double A[10000];
int th_id;
omp_set_num_threads(4);
#pragma omp parallel
{
    th_id = omp_get_thread_num();
    compute(th_id, A);
}
printf("Done.");
```

th\_id is shared.

### Remember the section(s)?

```
int N = 10000;
int data[N];
init_from_file( data ); // Initialize the data
#pragma omp parallel num_threads(4)
{
    #pragma omp sections
    {
        #pragma omp section
        compute(0,N,data) ;
        #pragma omp section
        compute(1,N,data) ;
        #pragma omp section
        compute(2,N,data) ;
        #pragma omp section
        compute(3,N,data) ;
    }
    print_vector( data );
}
```

```
compute(int index, int N, int* data) {
    int i_start, i_end, i ;

    i_start = index*N/4+1;
    i_end = (index+1)*N/4;
    printf("Working on data[%d, %d]\n",
           i_start, i_end);
    for (i=i_start; i<i_end;i++)
        data[i] = i*i
}
```


## Remember the section(s)?

```
int N = 10000;
int data[N];
init_from_file( data ); // Initialize the data
#pragma omp parallel num_threads(4)
{
    #pragma omp sections
    {
        #pragma omp section
        compute(0,N,data);
        #pragma omp section
        compute(1,N,data);
        #pragma omp section
        compute(2,N,data);
        #pragma omp section
        compute(3,N,data);
    }
}
print_vector( data );
```

```
compute(int index, int N, int* data) {
    int i_start, i_end, i;
    i_start = index*N/4+1;
    i_end = (index+1)*N/4;
    printf("Working on data[%d, %d]\n",
           i_start, i_end);
    for (i=i_start; i<i_end;i++)
        data[i] = i*i;
}
```

## Agenda

- More about the control of the concurrent accesses to the variables.
- How do you schedule the iterations of a loop among the threads?



## compute( ) reloaded

```
compute(int index, int N, int* data) {
    int i_start, i_end, i;
    int norm;
    i_start = index*N/4+1;
    i_end = (index+1)*N/4;
    printf("Working on data[%d, %d]\n",
           i_start, i_end);
    for (i=i_start; i<i_end;i++) {
        norm += data[i]*data[i];
    }
}
```

Is there a problem?

## Remember the section(s)?

```
int N = 10000;
int data[N]; int norm = 0;
init_from_file( data ); // Initialize the data
#pragma omp parallel num_threads(4) private(norm)
{
    #pragma omp sections
    {
        #pragma omp section
        compute(0,N,data);
        #pragma omp section
        compute(1,N,data);
        #pragma omp section
        compute(2,N,data);
        #pragma omp section
        compute(3,N,data);
    }
}
print_vector( data );
```

```
compute(int index, int N, int* data) {
    int i_start, i_end, i;
    i_start = index*N/4+1;
    i_end = (index+1)*N/4;
    printf("Working on data[%d, %d]\n",
           i_start, i_end);
    for (i=i_start; i<i_end;i++)
        norm += data[i]*data[i];
}
```

## This will not work!

- Each thread owns a **private copy** of `norm`.
- Each copy is free'd when the threads finish.
- The **global var** `norm` that remains when only the master thread keeps running **does not know anything** about the former local copies.



## Firstprivate/lastprivate

- The pragma **FIRSTPRIVATE(var1)** creates private versions of var1, copying the initial value of the global version.
- The pragma **LASTPRIVATE(var1)** creates private versions of var1, copying the final value(\*) to the global version.
  - (\*) as defined by the sequential execution.

### Remember the section(s)?

```
int N = 10000;
int data[N]; int norm = 0;
init_from_file( data ); // Initialize the data
#pragma omp parallel num_threads(4) lastprivate(norm)
{
    #pragma omp sections
    {
        #pragma omp section
        compute(0,N,data);
        #pragma omp section
        compute(1,N,data);
        #pragma omp section
        compute(2,N,data);
        #pragma omp section
        compute(3,N,data);
    }
    print_vector( data );
}
```

```
compute(int index, int N, int* data) {
    int i_start, i_end, i;

    i_start = index*N/4+1;
    i_end = (index+1)*N/4;
    printf("Working on data[%d, %d]\n",
           i_start, i_end);
    for (i=i_start; i<i_end; i++)
        norm += data[i]*data[i];
}
```

### The final problem

- This still does not work....
  - You do not want the last (partial) local value of the norm,
  - What you want is the **sum** of all the (partial) local values.
- This is called a **reduction**.
  - '+' is the reduction operator.
  - You can also use \*, or, and, MAX, MIN.
  - OpenMP syntax: **reduction(+:norm)**

### Remember the section(s)?

```
int N = 10000;
int data[N]; int norm = 0;
init_from_file( data ); // Initialize the data
#pragma omp parallel num_threads(4) reduction(+:norm)
{
    #pragma omp sections
    {
        #pragma omp section
        compute(0,N,data);
        #pragma omp section
        compute(1,N,data);
        #pragma omp section
        compute(2,N,data);
        #pragma omp section
        compute(3,N,data);
    }
    print_vector( data );
}
```

```
compute(int index, int N, int* data) {
    int i_start, i_end, i;

    i_start = index*N/4+1;
    i_end = (index+1)*N/4;
    printf("Working on data[%d, %d]\n",
           i_start, i_end);
    for (i=i_start; i<i_end; i++)
        norm += data[i]*data[i];
}
```

### Parallel for

```
int N = 10000;
int data[N]; int norm = 0;
init_from_file( data ); // Initialize the data
#pragma omp parallel num_threads(4) reduction
(+:norm)
{
    #pragma omp for
    for (i=0; i<N; i++)
        norm += data[i]*data[i];
}
print_vector( data );
```

← This pragma automatically distributes the iterations of the loop among the threads

### Parallel for

```
int N = 10000;
int data[N]; int norm = 0;
init_from_file( data ); // Initialize the data
#pragma omp parallel num_threads(4) reduction
(+:norm)
{
    #pragma omp for
    for (i=0; i<N; i++)
        norm += data[i]*data[i];
}
print_vector( data );
```

← Beware the limits of the loop!  
No write conflict!

### Parallel for – short version

```
int N = 10000;
int data[N]; int norm = 0;
init_from_file( data ); // Initialize the data
#pragma omp parallel for num_threads(4) reduction
(+:norm)
for (i=0; i<N; i++)
    norm += data[i]*data[i];
print_vector( data );
```

## Nested loops

- Let us consider a matrix product
  - 3 nested loops i, j, k.
- The **2 external loops** (i, j) can be run in parallel.
  - Actually, the 3rd one also (Reduce)
- So where do you put the “parallel for”?

## Options

```
for (i=0 ; i<N ; i++) {
    for (j=0 ; j<N ; j++) {
        C[i][j] = 0;

        for (k=0 ; k<N; k++)
            C[i][j] += A[i][k]*B[k][j];
    }
}
```

## Option 1

```
#pragma omp parallel for
for (i=0 ; i<N ; i++) {
    for (j=0 ; j<N ; j++) {
        C[i][j] = 0;

        for (k=0 ; k<N; k++)
            C[i][j] += A[i][k]*B[k][j];
    }
}
```

## Option 1

```
#pragma omp parallel for private(j,k)
for (i=0 ; i<N ; i++) {
    for (j=0 ; j<N ; j++) {
        C[i][j] = 0;

        for (k=0 ; k<N; k++)
            C[i][j] += A[i][k]*B[k][j];
    }
}
```

Beware the  
Private!

## Option 2

```
for (i=0 ; i<N ; i++) {
    #pragma omp parallel for private(k)
    for (j=0 ; j<N ; j++) {
        C[i][j] = 0;


        for (k=0 ; k<N; k++)
            C[i][j] += A[i][k]*B[k][j];
    }
}
```

## Option 2

```
for (i=0 ; i<N ; i++) {
    #pragma omp parallel for private(k)
    for (j=0 ; j<N ; j++) {
        C[i][j] = 0;

        for (k=0 ; k<N; k++)
            C[i][j] += A[i][k]*B[k][j];
    }
}
```

Beware the  
synch!



## Option 3

```

for (i=0 ; i<N ; i++) {
  #pragma omp parallel for private(k)
  for (j=0 ; j<N ; j++) {
    double tmp= 0;
    #pragma omp parallel for reduction(+:tmp)
    for (k=0 ; k<N; k++)
      tmp += A[i][k]*B[k][j];
    C[i][j] = tmp;
  }
}

```

## Options 1+2+3

```

#pragma omp parallel for private(j,k)
for (i=0 ; i<N ; i++) {
  #pragma omp parallel for
  for (j=0 ; j<N ; j++) {
    double tmp= 0;
    #pragma omp parallel for reduction(+:tmp)
    for (k=0 ; k<N; k++)
      tmp += A[i][k]*B[k][j];
    C[i][j] = tmp;
  }
}

```


## Distribution of the iterations

- How do you know **which thread runs which** values (between 0 and N-1) of the loop?
  - With sections the mapping was explicitly computed by the programmer!
- Do you really want to know?
  - OpenMP does it for you.

## Small example

N = 14 iterations  
P = 3 threads

## Small example



Each thread runs a **block** of approx.  $B=N/p$  +1 **contiguous iterations**:

- the 1st  $N/p$  threads run B iterations,
- the last threads run B-1 iterations

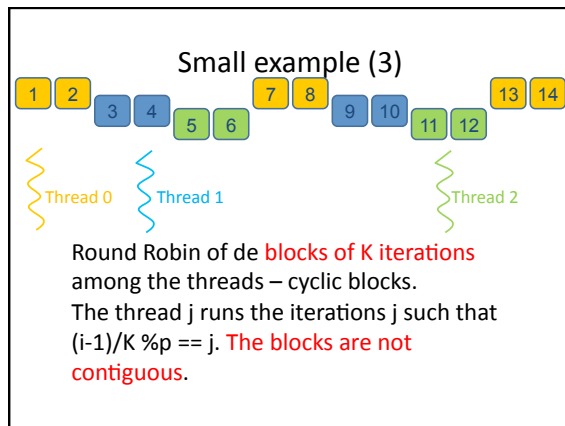
## Small example (2)



The allocation of iterations is **cyclic** (Round-Robin).

The thread j runs the iterations i such that  $i-1\%p == j$ . **They are not contiguous**.

Each thread runs approx. **B iterações**.



### Pros / Cons

Grouping the iterations in block increases the locality.

Having small groups balances the load.

You want to increase K

You want to decrease K

### Controlling the mapping with OpenMP

- You use `schedule(static, K)`, together with “parallel for”:

```
#pragma omp parallel for schedule(static, 5)
for (i=0; i<N; i++) { //...
```

- The first parameter specifies the schedule.
  - static, dynamic, guided
- The second parameter sets the **size** of the block.

### Conclusion: using `schedule()`


- You always should use “Schedule”.
- You should experiment to find the right value for K.
- If the computation is **homogeneous**, then `schedule(static, N/p)` should do fine.
  - This is the default.

### Próxima aula

### Extra stuff

## Paralelismo de Laços

- `#pragma omp parallel for`
- `for (i = 0; i < n; i++){`
- `c[i] = a[i] + b[i];`
- `}`



- Demonstração: exemplo de paralelismo de laços
- Restrições de acesso:

```
int fator = 0.5;
#pragma omp parallel for private( fator )
for ( i = 0; i < N; i++){
    c[i] = fator * a[i];
}
```

```
#pragma omp parallel for private( j );
for ( i = 0; i < N; i++){
    for( j = 0; j < N; j++){
        c[i] = a[i] + b[j];
    }
}
```

37

\*firstprivate, lastprivate