

Parallel Programming


Trapeze method with MPI

Henrique Weber
João Luiz Grave Gross


Professor Nicolas Mailard

December 3rd, 2010.


Area under a function



Area under a function




Trapeze method



$$A = \sum_{i=0}^{n-1} A_i$$

$$f(x) = \left(\frac{(b - a)}{n} \right) / 2 * (f_0 + f_n + \sum_{j=1}^{n-1} f_j)$$


Trapeze method



$$A = \sum_{i=0}^{n-1} A_i$$

$$f(x) = \left(\frac{(b - a)}{n} \right) \cdot \left(\frac{f_0 + f_n}{2} + \sum_{j=1}^{n-1} f_j \right)$$

Trapeze method



$$A = \sum_{i=0}^{n-1} A_i$$

$$f(x) = \left(\frac{(b - a)}{n} \right) / 2 * (f_0 + f_n + \sum f_j, \text{ from } j=1 \text{ to } n-1)$$

Sequential Program

Functions

```
int main ()
{
    /* Variable Declaration */

    /* Time Counting Begin */
    CalculateInteger( interval, polynomial, &subintervals, &integer, &polynomial_grade );
    /* Time Counting End */

    PrintResult( interval, polynomial, &subintervals, &integer, &t_total, &polynomial_grade );

    return 0;
}
```

Simplified main function of the program.

Sequential Program

Functions

- 1) **double** * CaptureEntries(**double** *interval, **long int** *subintervals, **long int** *polynomial_grade)
- 2) **void** CalculateInteger(**double** *interval, **double** *polynomial, **long int** *subintervals, **double** *integer, **long int** *polynomial_grade)
- 3) **double** CalculatePolynomial(**double** *polynomial, **double** *xi, *polynomial_grade)
- 4) **void** PrintResult(**double** *interval, **double** *polynomial, **long int** *subintervals, **double** *integer, **double** *t_total, **long int** *polynomial_grade)

Sequential Program

Functions

1) **double** * CaptureEntries(**double** *interval, **long int** *subintervals, **long int** *polynomial_grade)

This function just capture all the entries needed for the execution of the trapeze method.

Sequential Program

Functions

```
2) void CalculateInteger( double *interval, double *polynomial, long int *subintervals,
double *integer, long int *polynomial_grade )
{
    double interval_h, interval_aux;
    long int i;

    interval_h = (interval[1] - interval[0]) / *subintervals;
    *integer = *integer + CalculatePolynomial( polynomial, &interval[0], polynomial_grade );
    *integer = *integer + CalculatePolynomial( polynomial, &interval[1], polynomial_grade );

    for(i = 1; i < *subintervals; i++){
        interval_aux = interval[0] + i*interval_h;
        *integer = *integer + 2*CalculatePolynomial( polynomial, &interval_aux,
polynomial_grade );
    }
    *integer = *integer * interval_h / 2;
    return;
}
```

Sequential Program

Functions

3) **double** CalculatePolynomial(**double** *polynomial, **double** *xi, **long int** *polynomial_grade)

```
{  
    double trapeze = 0;  
    int i;  
  
    for(i = 0; i <= *polynomial_grade; i++)  
        trapeze = trapeze + polynomial[i]*pow(*xi,i);  
  
    return trapeze;  
}
```

Sequential Program

Functions

4) **void** PrintResult(**double** *interval, **double** *polynomial, **long int** *subintervals, **double** *integer, **double** *t_total, **long int** *polynomial_grade)

Just prints the result of the program.

Sequential Program

Time Measurement

```
...
#include <sys/time.h>

int main()
{
    ...
    struct timeval tv_begin, tv_end;
    double t_begin, t_end, t_total;

    ...
    gettimeofday(&tv_begin, NULL);
    t_begin = (tv_begin.tv_sec)*1000000 + tv_begin.tv_usec;
    CalculateInteger( interval, polynomial, &subintervals, &integer, &polynomial_grade );
    gettimeofday(&tv_end, NULL);
    t_end = (tv_end.tv_sec)*1000000 + tv_end.tv_usec;

    t_total = t_end - t_begin; //The value is in micro seconds
    ...
}
```

Sequential Program

Execution Enviroment:

two quad-core Intel Xeon E5530 (Nehalem arch)
64KB L1 cache, 12MB L2 cache, 2GHz

How to compile:

```
gcc trapeze_sequential.c -o trapeze_sequential
```

How to run:

```
./trapeze_sequential
```

Input: Polynomial: $0.625x^4 - 4x^3 + 2x + 1$
Interval: $[0,8]$
Subintervals: 10^7

Output: 72,00000

Average Time = 2.841356 s

Parallel Program

Differences from the sequential version

- 1) More than one process in the program execution
- 2) One root process
- 3) n slave processes
- 4) Task division
- 5) MPI functions

Similarities from the sequential version

- 1) Same algorithm
- 2) Same functions (just adapted)
- 3) Just one program written (just adapted for parallel uses)

Parallel Program

Functions

```
...
#include <mpi.h>

int main ()
{
    /* Variable declaration */
    /* MPI initializing */

    if (rank == 0){ /* Root process */
        /* Capture entries */

        /* Time counting begin */
        /* Divide tasks */
        /* Receive and sum results */
        /* Time counting end */
    }
    ...
}
```


Parallel Program

Functions

```
...  
else{ /* Slave process */  
    /* Receive information */  
    /* Compute part of the problem */  
    /* Send results to root */  
}  
  
/* MPI ending */  
return 0;  
}
```

Parallel Program

Functions

1) **double** * CaptureEntries(**double** *interval, **long int** *subintervals, **long int** *polynomial_grade)

2) **void** CalculateInteger(**int** rank, **int** numtasks, **double** interval_h, **double** *interval, **double** *polynomial, **long int** *subintervals, **double** *parcial_integer, **long int** *polynomial_grade, **long int** init_subint, **long int** end_subint)

3) **double** CalculatePolynomial(**double** *polynomial, **double** *xi, *polynomial_grade)

4) **void** PrintResult(**double** *interval, **double** *polynomial, **long int** *subintervals, **double** *integer, **double** *t_total, **long int** *polynomial_grade)

All the 4 functions works just like the functions from the sequential code.

Parallel Program

Functions

```
2) void CalculateInteger(int rank, int numtasks, double interval_h, double *interval,
double *polynomial, long int *subintervals, double *parcial_integer, long int *polynomial_grade,
long int init_subint, long int end_subint )
{
    double interval_aux;
    long int i;
    *parcial_integer = 0;

    for(i = init_subint; i < end_subint; i++){
        interval_aux = interval[0] + i*interval_h;
        *parcial_integer = *parcial_integer + 2*CalculatePolynomial( polynomial,
&interval_aux, polynomial_grade );
    }

    return;
}
```

Parallel Program

Time Measurement

The same as in the sequential code.

Implemented in the root process.

Parallel Program

Data transfer through processes

```
int main ()
{
    ...
    if (rank == 0){ /* Root process */
        ...
        /* Time counting begin */

        MPI_Bcast( &polynomial_grade, 1, MPI_LONG, root, MPI_COMM_WORLD );
        /* Broacast other variables */

        ...

        for(i = 1; i < numtasks; i++){
            MPI_Send(&init_subint, 1, MPI_DOUBLE, i, tag, MPI_COMM_WORLD);
            MPI_Send(&end_subint, 1, MPI_DOUBLE, i, tag, MPI_COMM_WORLD);
            ...
        }

        ...
    }
}
```

Parallel Program

Data transfer through processes

```
...
for(i = 1; i < numtasks; i++){
    ...
    MPI_Recv(&parcial_integer, 1, MPI_DOUBLE, i, tag, MPI_COMM_WORLD,
&Stat);
    integer += parcial_integer;
}

/* Calculate final integer */
/* Finish time counting */
}

...
```

Parallel Program

Data transfer through processes

```
else{ /* Slave process */
    MPI_Bcast( &polynomial_grade, 1, MPI_LONG, root, MPI_COMM_WORLD );
    /* Receive other data via broadcast */
    ...

    MPI_Recv(&init_subint, 1, MPI_DOUBLE, root, tag, MPI_COMM_WORLD, &Stat);

    MPI_Recv(&end_subint, 1, MPI_DOUBLE, root, tag, MPI_COMM_WORLD, &Stat);

    CalculateInteger( rank, numtasks, interval_h, interval, polynomial, &subintervals,
    &parcial_integer, &polynomial_grade, init_subint, end_subint );

    MPI_Send(&parcial_integer, 1, MPI_DOUBLE, dest, tag, MPI_COMM_WORLD);
}

MPI_Finalize();

return 0;
}
```

Parallel Program

How to compile:

`mpicc trapeze_parallel.c -o trapeze_parallel`

How to run:

`mpirun -np 2 ./trapeze_parallel` (here the number 2 is just an example, but it can be higher than 2)

Input: Polynomial: $0.625x^4 - 4x^3 + 2x + 1$

Interval: $[0,8]$

Subintervals: 10^7

Output: 72,00000

Number of processes	Average time (sec)	Number of processes	Average time (sec)
2	2.58	13	0.38
3	1.28	14	0.37
4	0.85	15	0.38
5	0.65	16	0.36
6	0.55	17	0.38
7	0.47	18	0.37
8	0.4	19	0.38
9	0.38	20	0.38
10	0.43	21	0.38
11	0.41	22	0.39
12	0.4		

Parallel Program

How to compile:

`mpicc trapeze_parallel.c -o trapeze_parallel`

How to run:


`mpirun -np 2 ./trapeze_parallel` (here the number 2 is just an example, but it can be higher than 2)

Input: Polynomial: $0.625x^4 - 4x^3 + 2x + 1$

Interval: $[0,8]$

Subintervals: 10^7

Output: 72,00000



Sequential Program x Parallel algorithm

Sequential Algorithm	Parallel Algorithm	Gain
2.841356 sec	Worst case (2.58 sec) – 2 processes	9,19%
	Medium case (0.55 sec)	516%
	Better case (0.36 sec) – 16 processes	780%

Conclusion

- Parallel applications can provide a huge gain in performance.
- Problems can be fully or partially parallelized, we just have to SEE and THINK how this can be done.
- Parallel programming is growing fast so is important for us to master it.
- Its even more important due to the increase of cores inside processors, in other words, parallel application will be seen more frequently.

Thanks!

Questions?!

Henrique Weber - hweber@inf.ufrgs.br

João Gross - jlggross@inf.ufrgs.br